

Forwarding and Routing

Richard T. B. Ma

School of Computing

National University of Singapore

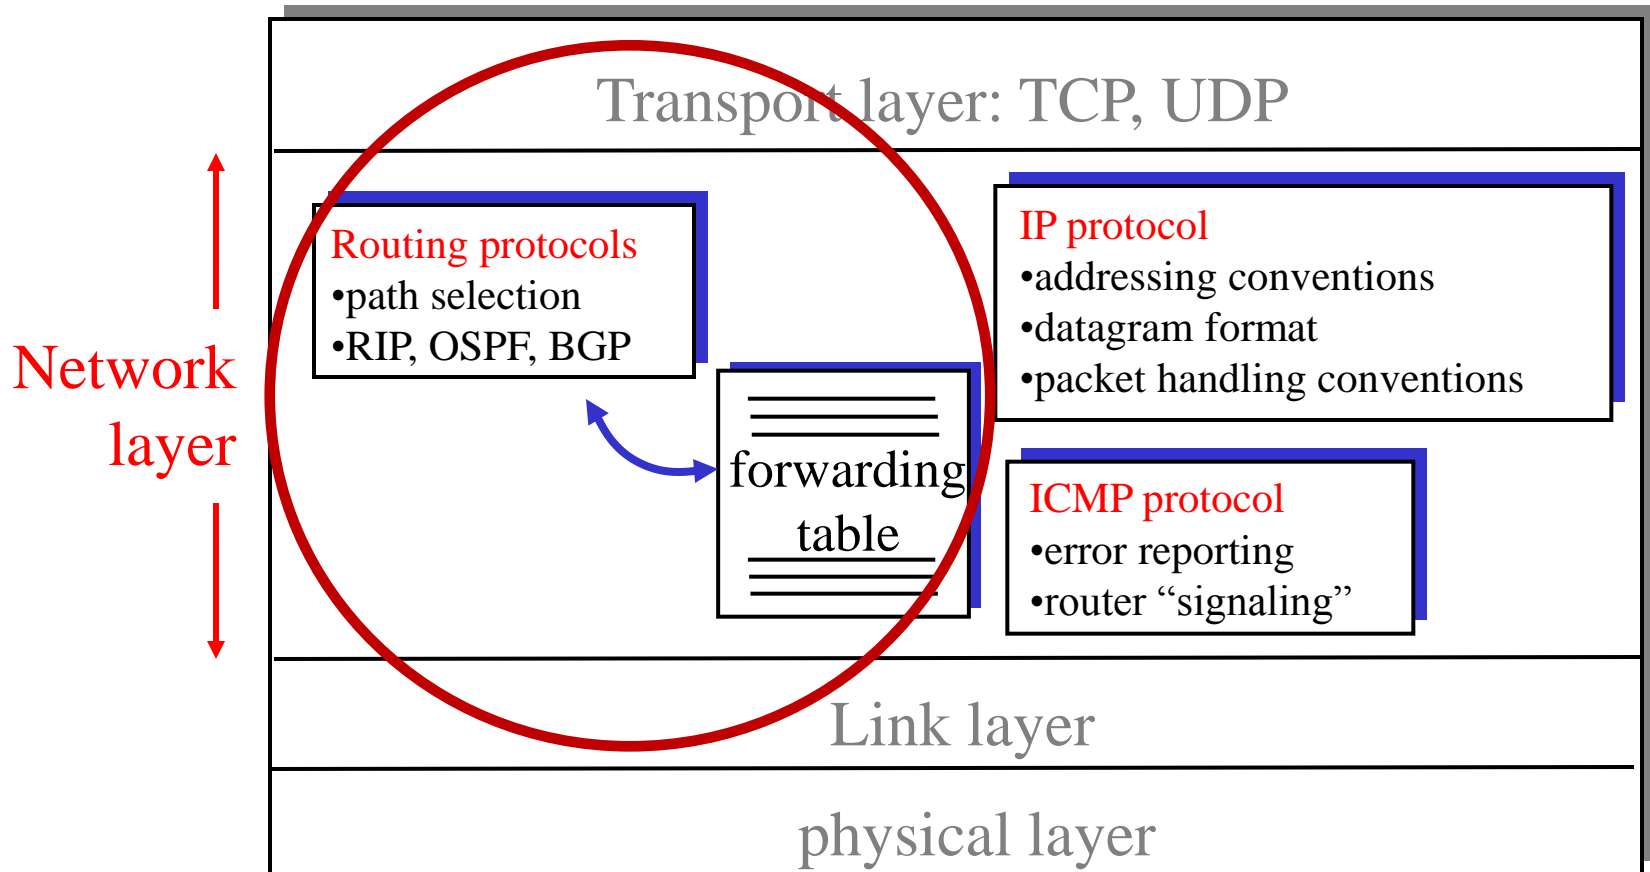
CS 3103: Compute Networks and Protocols

IP Protocol Stack: Key Abstractions

Application	Applications	
Transport	Reliable streams	Unreliable datagrams
Network	Best-effort <i>global</i> packet delivery	
Link	Best-effort <i>local</i> packet delivery	


The Internet Network layer

Host, router network layer functions:




Routing in virtual circuits

- ❑ used to setup, maintain teardown VC
- ❑ used in ATM, frame-relay, X.25
- ❑ not used in today's Internet



Datagram networks

- ❑ no call setup at network layer
- ❑ routers: no end-to-end state information
 - ❖ no network-level concept of “connection”
- ❑ packets forwarded using destination host address



Two Key Network-Layer Functions

□ *Forwarding (data plane):*

move packets from an input interface to an appropriate output interface in a router

□ *Routing (control plane):*

determine route taken by packets from source to destination

❖ *routing algorithms*

analogy:

❖ **routing:** process of planning trip from source to destination

❖ **forwarding:** process of getting through single interchange

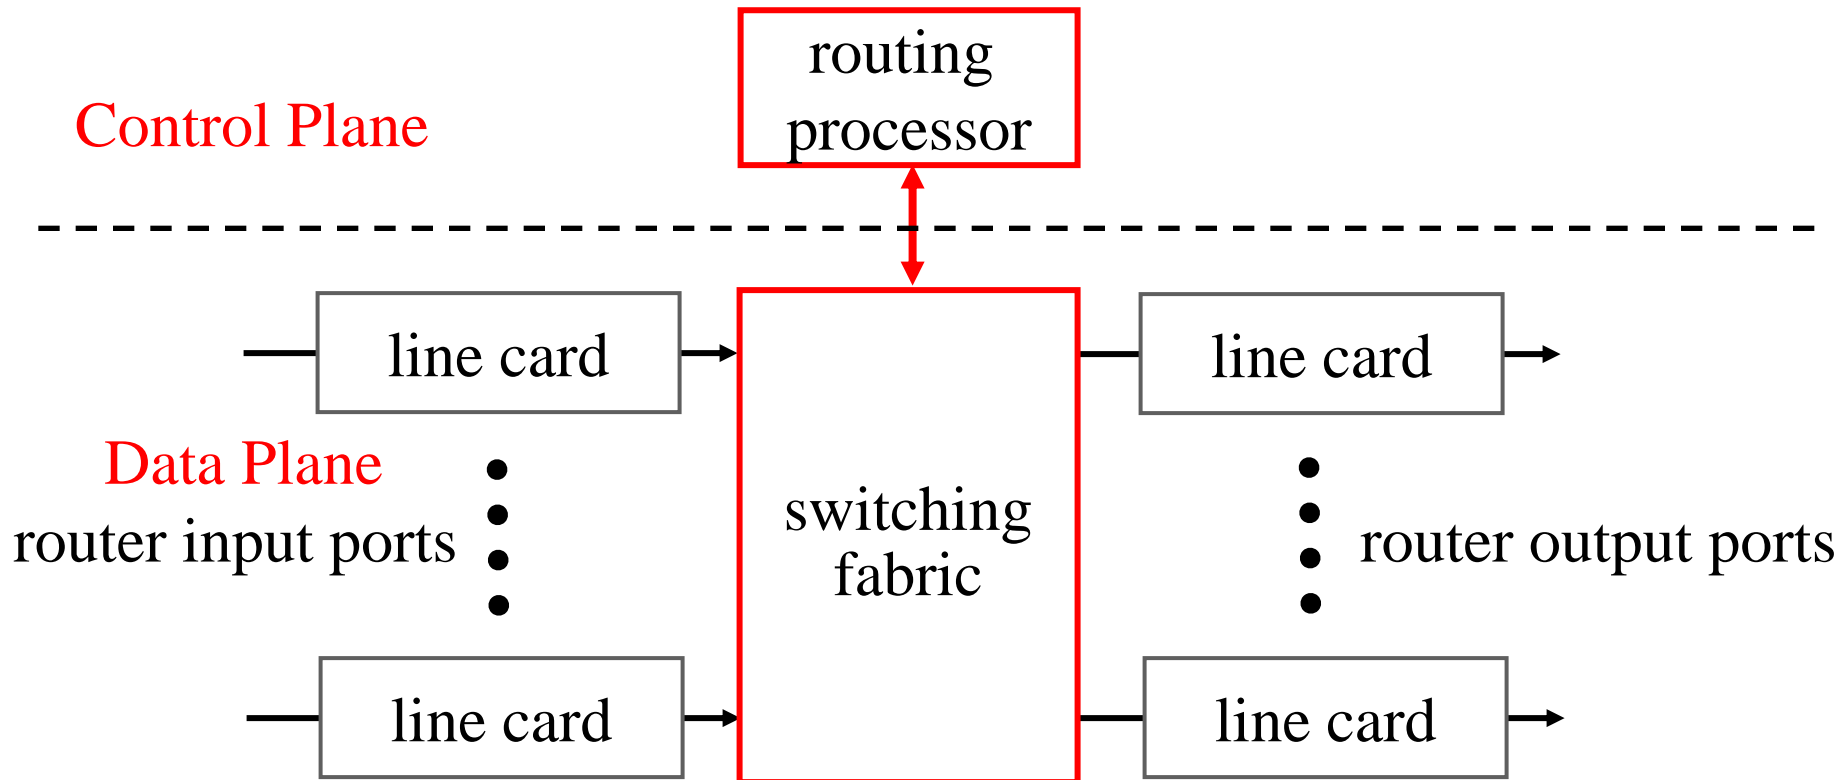
Forwarding: Hop-by-Hop

- ❑ Each router has a forwarding table
 - ❖ maps destination addresses to outgoing interfaces
- ❑ Upon receiving a packet
 - ❖ inspect the destination IP address in the header index into the table
 - ❖ determine the outgoing interface
 - ❖ forward the packet out that interface
- ❑ Then, the next router on the path repeats
 - ❖ and the packet travels along the path to the destination

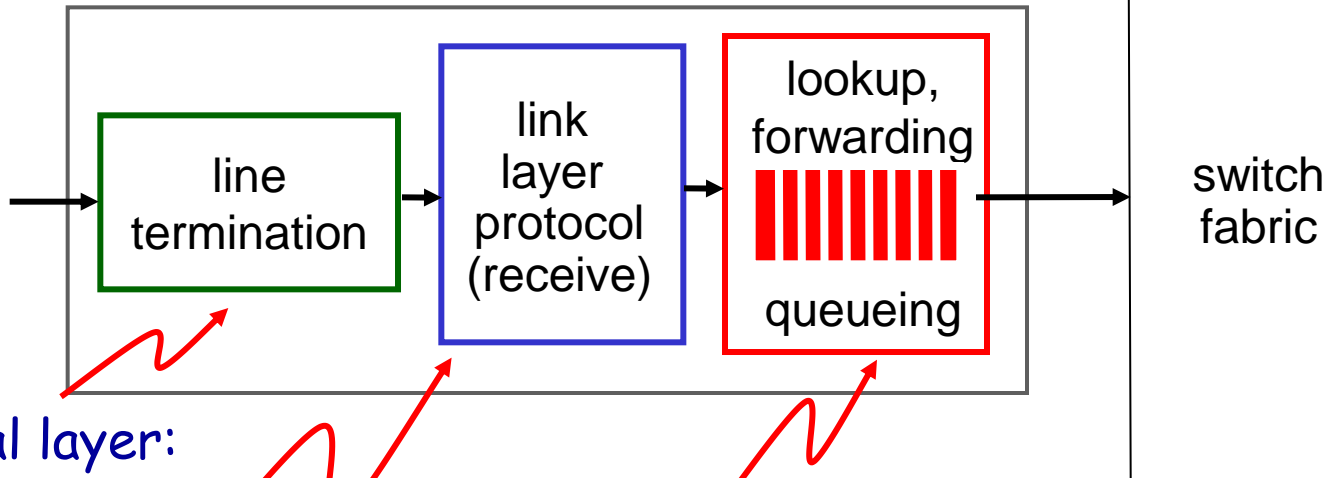
Router Architecture Overview

two key router functions:

- ❑ run routing algorithms/protocol (RIP, OSPF, BGP)
- ❑ *forwarding* datagrams from incoming to outgoing link



Input port functions



physical layer:
bit-level reception


data link layer:
e.g., Ethernet

decentralized switching:

- ❑ given datagram destination, lookup output port using forwarding table in input port memory (*"match plus action"*)
- ❑ goal: complete input port processing at "line speed"
- ❑ queuing: if datagrams arrive faster than forwarding rate into switch fabric

Switching fabrics


- ❖ transfer packet from input buffer to appropriate output buffer
- ❖ switching rate: rate at which packets can be transfer from inputs to outputs
 - often measured as multiple of input/output line rate
 - N inputs: switching rate N times line rate desirable
- ❖ three types of switching fabrics



Switching via memory


first generation routers:

- ❑ traditional computers with switching under direct control of CPU
- ❑ packet copied to system's memory
- ❑ speed limited by memory bandwidth (2 bus crossings per datagram)




Switching via a bus

- ❖ datagram from input port memory to output port memory via a shared bus
- ❖ internal label is used to indicate output port
- ❖ *bus contention*: switching speed limited by bus bandwidth
- ❖ 32 Gbps bus, Cisco 5600: sufficient speed for access and enterprise routers






Switching via interconnection nets

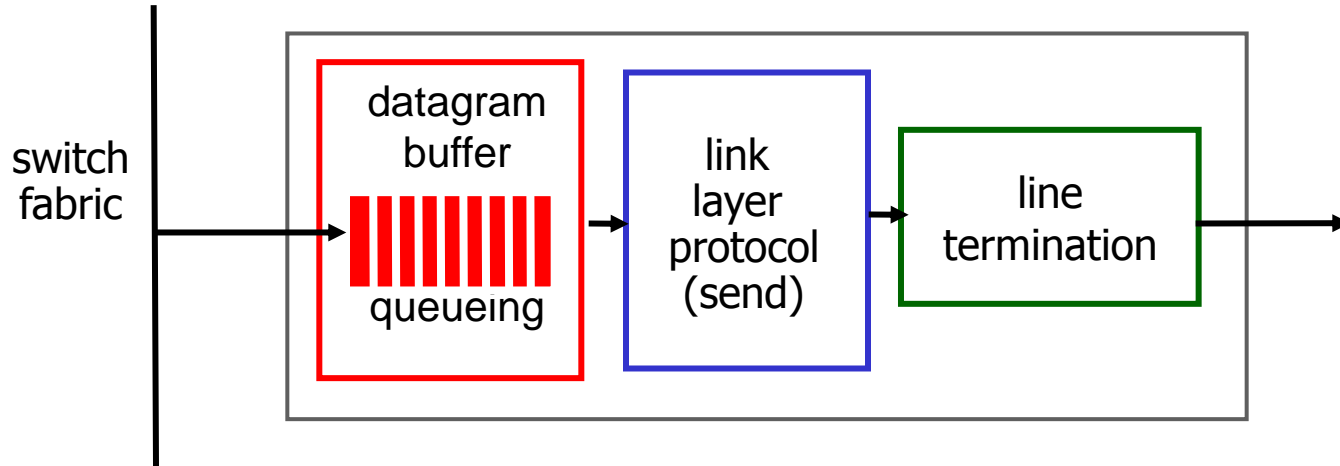
- ❖ overcome bus bandwidth limitations
- ❖ banyan networks, crossbar, other interconnection nets initially developed to connect processors in multiprocessor
- ❖ advanced design: fragmenting datagram into fixed length cells, switch cells through the fabric.
- ❖ Cisco 12000: switches 60 Gbps through the interconnection network



Banyan networks (*)

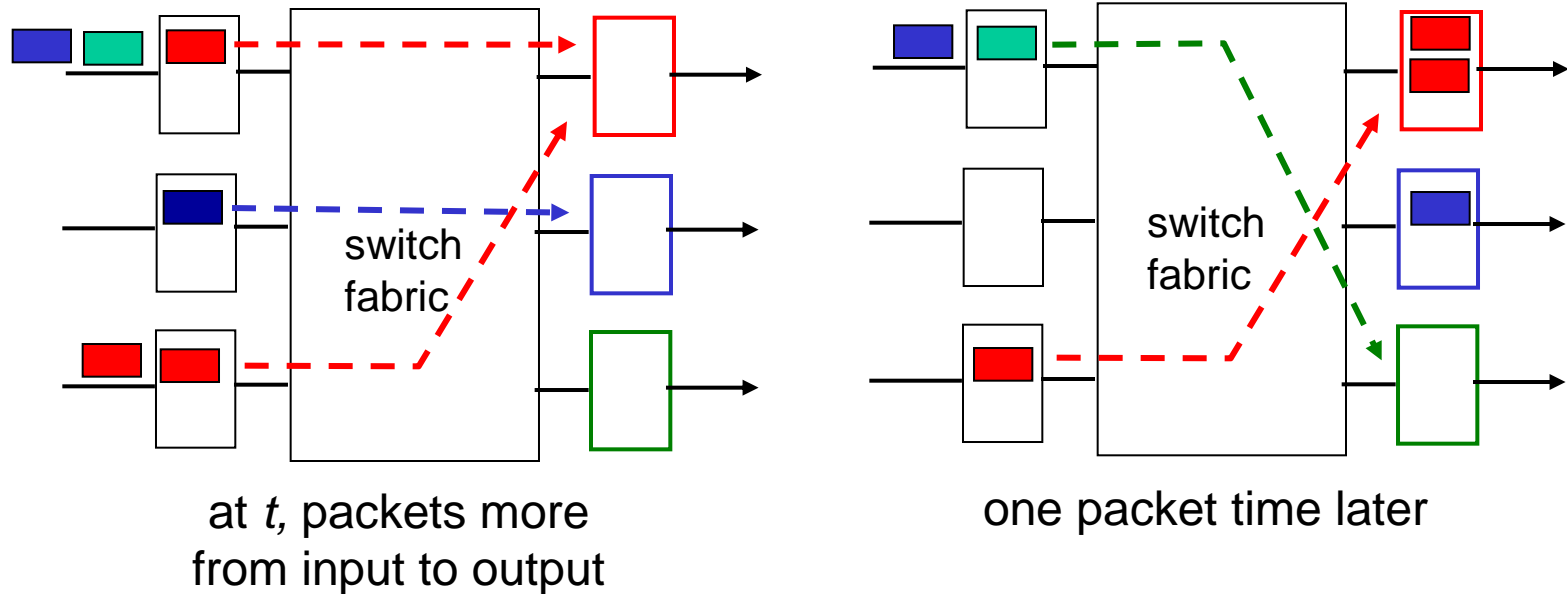


Output ports



- ❖ *buffering* required when datagrams arrive from fabric faster than the transmission rate
- ❖ *scheduling discipline* chooses among queued datagrams for transmission


Output port queueing




- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to congestion and output port buffer overflow!*

Input port queuing

- ❑ fabric slower than input ports combined -> queueing may occur at input queues
 - ❖ *delay and loss due to input buffer overflow!*
- ❑ **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward




**output port contention:
only one red packet can be transferred
(lower red packet is blocked)**



**one packet time later:
green packet experiences
HOL blocking**

Routing Processor



- ❑ “loopback” interface
 - ❖ IP address of the CPU on the router
- ❑ “Control-Plane” software
 - ❖ implementation of routing protocols
 - ❖ creation of forwarding table for the line cards
- ❑ Handling of special data packets
 - ❖ packets with IP options enabled
 - ❖ packets with expired Time-to-Live
- ❑ Network management functions
 - ❖ command-line interface (CLI) for configuration
 - ❖ Transmission of measurement statistics

Routing vs. Forwarding

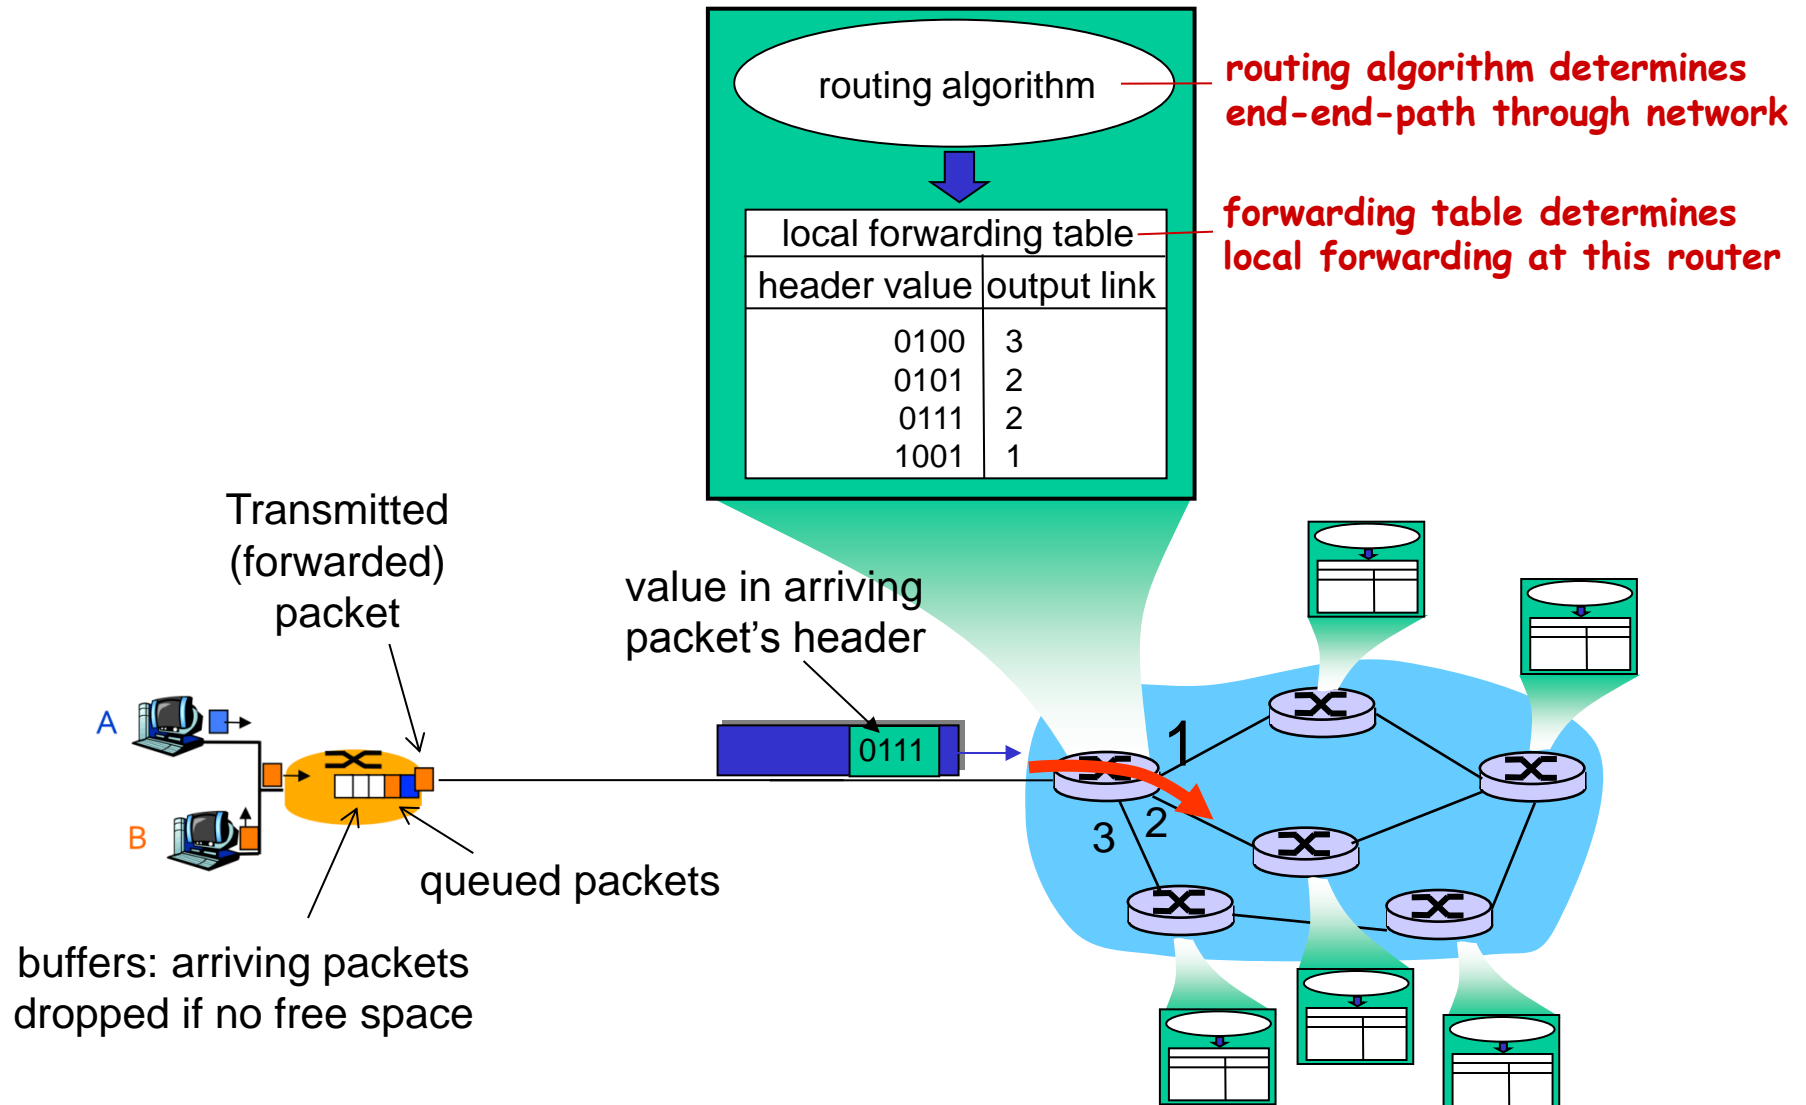
□ Routing: control plane

- ❖ Computing paths the packets will follow
- ❖ Routers talking amongst themselves
- ❖ Creating the forwarding tables

□ Forwarding: data plane

- ❖ Directing a data packet to an outgoing link
- ❖ Using the forwarding tables

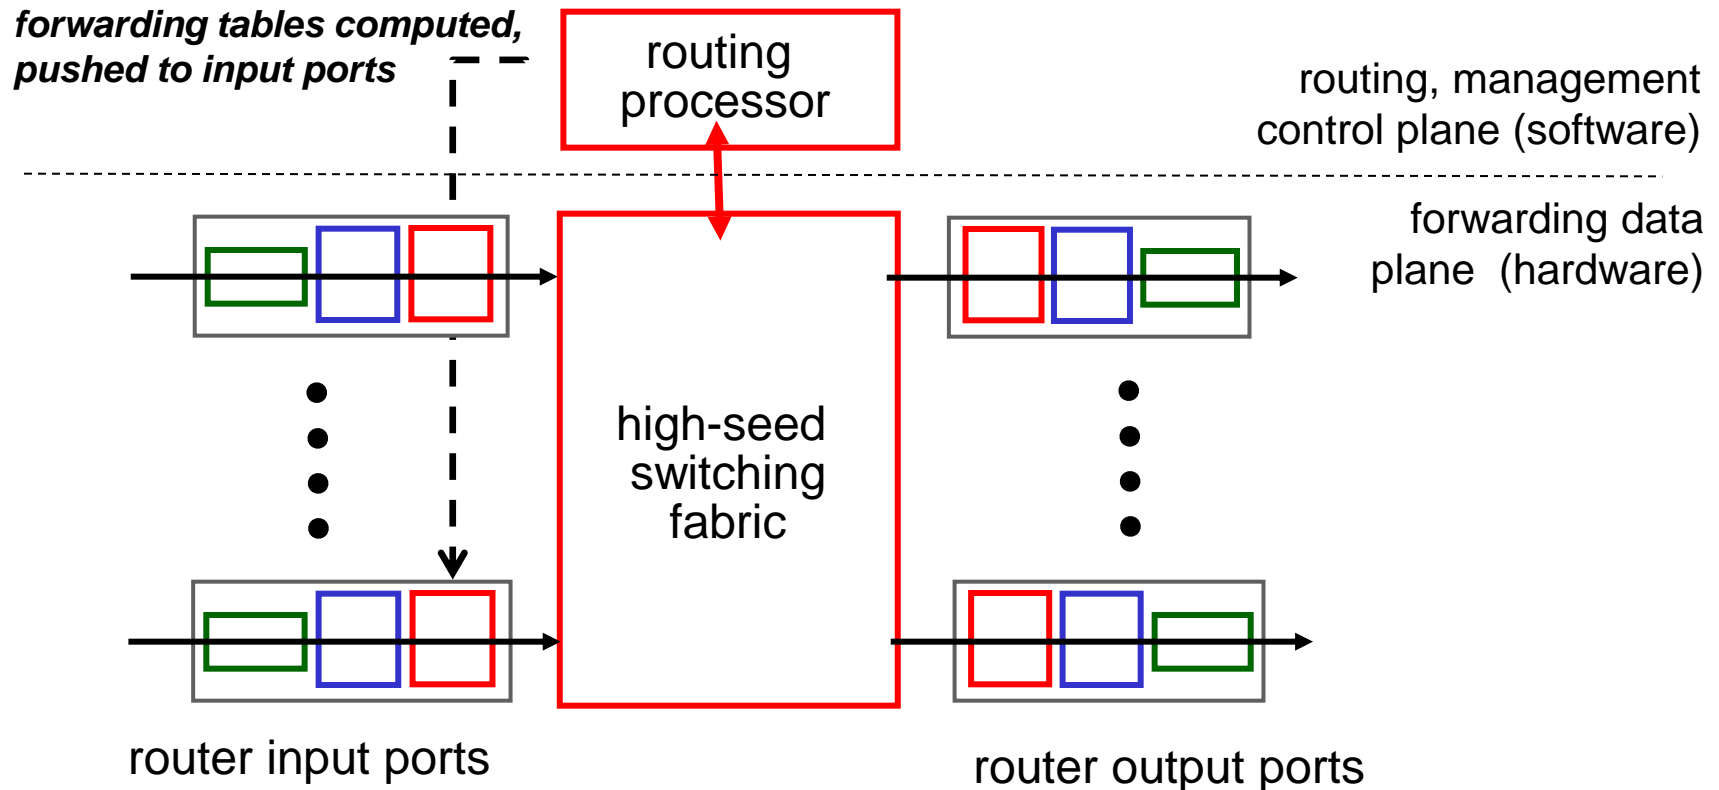
Interplay between routing and forwarding



Router architecture overview

two key router functions:

- ❖ run routing algorithms/protocol (RIP, OSPF, BGP)
- ❖ *forwarding* datagrams from incoming to outgoing link




Goal of routing

- ❑ You may be able to send datagrams directly to the destination
- ❑ If not, the router attempts to send datagrams to a router that is nearer the destination.
- ❑ The *goal* of a routing protocol is very simple:
Find a “good” path from source to destination.
 - ❖ Hosts often have default routers/gateways
 - ❖ We can focus on the routing from the source router to the destination router

Computing Paths between Routers

- ❑ Routers need to know two things
 - ❖ which router to use to reach a destination?
 - ❖ which interface to use to reach that router?




Routing in the Internet: Example


- ❑ How does a router construct its forwarding table?
- ❑ How does a router know which is the next hop towards a destination?
- ❑ Use a **routing protocol** to propagate (and update) reachability information

Router 1's table

Destination	Next Hop
2	2
3	3
4	3



Graph abstraction



Graph: $G = (N, E)$


N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Remark: Graph abstraction is useful in other network contexts

Example: P2P, where N is set of peers and E is set of TCP connections

Graph abstraction: costs



- $c(x, x') = \text{cost of link } (x, x')$

- e.g., $c(w, z) = 5$

- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Question: What's the least-cost path between u and z ?

Routing algorithm: algorithm that finds least-cost path

Routing algorithm classification

Q: global or decentralized information?

global:

- ❑ all routers have complete topology, link cost info
- ❑ “link state” algorithms

decentralized:

- ❑ router knows physically-connected neighbors, link costs to neighbors
- ❑ iterative process of computation, exchange of info with neighbors
- ❑ “distance vector” algorithms

Q: static or dynamic?

static:

- ❖ routes change slowly over time

dynamic:

- ❖ routes change more quickly
 - periodic update
 - in response to link cost changes

A Link-State Routing Algorithm

Dijkstra's algorithm

- net topology, link costs known to all nodes
 - ❖ accomplished via “link state broadcast”
 - ❖ all nodes have same info
- computes least cost paths from one node (source) to all others
 - ❖ gives *forwarding table* for that node
- after k iterations, know least-cost path to k destinations

notation:

- $c(x,y)$: link cost from node x to y ; $= \infty$ if not direct neighbors
- $D(v)$: current value of cost of path from source to destination v
- $p(v)$: predecessor node along path from source to v
- N' : set of nodes whose least cost path definitively known

Dijkstra's Algorithm

1 **Initialization:**

2 $N' = \{u\}$

3 for all nodes v

4 if v adjacent to u

5 then $D(v) = c(u,v)$

6 else $D(v) = \infty$

7

8 **Loop**

9 find w not in N' such that $D(w)$ is a minimum

10 add w to N'

11 update $D(v)$ for all v adjacent to w and not in N' :

12 **$D(v) = \min(D(v), D(w) + c(w,v))$**

13 /* new cost to v is either old cost to v or known

14 shortest path cost to w plus cost from w to v */

15 **until all nodes in N'**




Dijsktra's algorithm: example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uwx	6,w			11,w	14,x
3	uwxv				10,v	14,x
4	uwxvy					12,y
5	uwxvyz					


notes:

- ❖ construct shortest path tree by tracing predecessor nodes
- ❖ ties can exist (can be broken arbitrarily)




Dijkstra's algorithm: example 2

Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					



Dijkstra's algorithm: example 2

resulting shortest-path tree from u:



resulting forwarding table in u:

destination	link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)


Dijkstra's algorithm, discussion

algorithm complexity: n nodes


- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

oscillations possible:


- ❖ e.g., suppose link cost = amount of carried traffic:




initially



**given these costs,
find new routing....
→ new costs**



**given these costs,
find new routing....
→ new costs**



**given these costs,
find new routing....
→ new costs**

Distance Vector (DV) algorithm

- ❑ distributed, iterative, and asynchronous
 - ❖ receives info from directly attached neighbors
 - ❖ continues on until no more info is exchanged
 - ❖ nodes can update and send info asynchronously
- ❑ $D_x(y)$ = estimate of least cost from x to y
 - ❖ x maintains distance vector $D_x = [D_x(y): y \in N]$
- ❑ For node x :
 - ❖ knows cost to each neighbor v : $c(x,v)$
 - ❖ also maintains its neighbors' DVs. For each neighbor v , x maintains $D_v = [D_v(y): y \in N]$

Distance Vector Algorithm

Bellman-Ford Equation (dynamic programming)

Define

$d_x(y) :=$ cost of least-cost path from x to y

Then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

cost to neighbor v

\uparrow


\uparrow

\min taken over all neighbors v of x

cost from neighbor v to destination y


The fact behind Bellman-Ford

❖ $d_x(y) = \min_v \{c(x,v) + d_v(y)\}$



$d_v(y)$	- shortest distance between v and y
$c(x,v)$	- cost between x and v
N	- number of neighbor of x

Bellman-Ford example



clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), c(u,x) + d_x(z), c(u,w) + d_w(z) \} \\ &= \min \{ \quad 2 + 5, \quad \quad 1 + 3, \quad \quad 5 + 3 \} = 4 \end{aligned}$$

❑ Practical contributions of B-F equation:

- ❖ node achieving minimum is next hop in the shortest path, used in forwarding table
- ❖ it suggests the form of the neighbor-to-neighbor communication used in the DV algorithm

Distance vector algorithm

key idea:

- ❖ from time-to-time, each node sends its own distance vector estimate to neighbors
- ❖ when x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{ c(x,v) + D_v(y) \} \text{ for each } y \in N$$

- ❖ under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $dx(y)$

Distance vector algorithm

iterative, asynchronous:

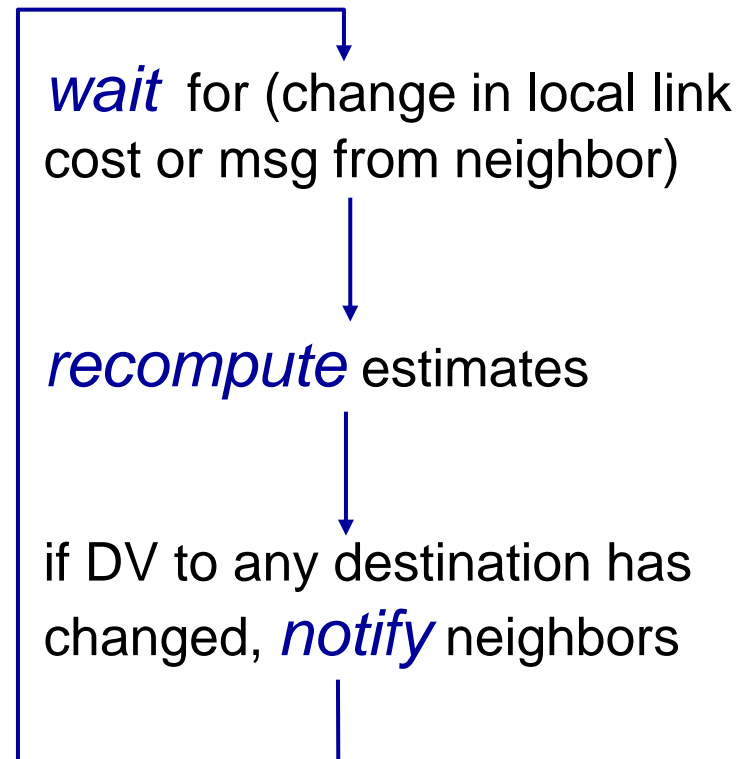
each local iteration caused by

- local link cost change
- DV update message from neighbor

distributed:

- each node notifies neighbors *only* when its DV changes
 - ❖ neighbors then notify their neighbors if necessary

at each node:



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\} \\ = \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\} \\ = \min\{2+1, 7+0\} = 3$$

**node x
table**

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

**node y
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

**node z
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0



time

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

**node x
table**

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

**node y
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

**node z
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0


		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0




time

Distance vector: link cost changes

link cost changes:

- ❖ node detects local link cost decrease
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors



“good
news
travels
fast”

t_0 : y detects link-cost change, updates its DV, informs its neighbors.


t_1 : z receives update from y, updates its table, computes new least cost to x, sends its neighbors its DV.

t_2 : y receives z's update, updates its distance table. y's least costs do *not* change, so y does *not* send a message to z.

Distance vector: link cost changes

link cost changes:

- ❖ node detects local link cost increase
- ❖ $D_y(x) = 4$, $D_y(z) = 1$, $D_z(y) = 1$, and $D_z(x) = 5$
- ❖ $D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\}$
 $= \min\{60 + 0, 1 + 5\} = 6$
- ❖ $D_z(x) = \min\{50 + 0, 1 + 6\} = 7$
- ❖ *bad news travels slow* - “count to infinity” problem!



poisoned reverse:

- ❖ If Z routes through Y to get to X :
 - Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)
- ❖ will this completely solve count to infinity problem?

Comparison of LS and DV algorithms

Message complexity

- ❑ LS: with n nodes, E links, $O(nE)$ msgs sent; talk to all nodes, only about direct links
- ❑ DV: exchange between neighbors only, but about least-distance to all nodes
 - ❖ convergence time varies

Speed of Convergence

- ❑ LS: $O(n^2)$ algorithm requires $O(nE)$ msgs
 - ❖ may have oscillations
- ❑ DV: convergence time varies
 - ❖ may be routing loops
 - ❖ count-to-infinity problem

Robustness: what if router malfunctions?

- ❑ LS: node can advertise incorrect *link* cost;
each computes only its *own* table
- ❑ DV: node can advertise incorrect *path* cost;
each table used by others (error propagate thru network)

Hierarchical routing

our routing study thus far - idealization

- ❖ all routers identical

- ❖ network “flat”

... *not* true in practice

scale: with 600 million destinations:

- ❑ can't store all dest's in routing tables!

- ❑ routing table exchange would swamp links!

administrative autonomy


- ❖ internet = network of networks

- ❖ each admin may want to control routing in its own network, or hide topology

Hierarchical routing

- aggregate routers into regions,
 - ❖ “autonomous systems” (AS)
- routers in same AS run same protocol
 - ❖ “intra-AS” routing protocol
 - ❖ routers in different AS can run different intra-AS routing protocol
- *gateway router*:
 - ❖ at “edge” of its own AS
 - ❖ has link to router in another AS

Interconnected ASes



- ❖ forwarding table configured by both intra- and inter-AS routing algorithm
 - intra-AS sets entries for internal dests
 - inter-AS & intra-AS sets entries for external dests


Inter-AS tasks

- ❖ suppose router in AS1 receives datagram destined outside of AS1
 - router should forward packet to gateway router, but which one?

AS1 must:


1. learn which destds are reachable through AS2, which through AS3
2. propagate this reachability info to all routers in AS1

job of inter-AS routing!




Set forwarding table in router 1d

- suppose AS1 learns (via inter-AS protocol) that subnet x reachable via AS3 (gateway 1c), but not via AS2
 - ❖ inter-AS protocol propagates reachability info to all internal routers
- router 1d determines from intra-AS routing info that its interface I is on the least cost path to 1c
 - ❖ installs forwarding table entry (x, I)



Example: choosing among multiple ASes

- now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- to configure forwarding table, router 1d must determine which gateway it should forward packets towards for dest **x**
 - ❖ this is also job of inter-AS routing protocol!



Example: choosing among multiple ASes

- now suppose AS1 learns from inter-AS protocol that subnet *x* is reachable from AS3 *and* from AS2.
- to configure forwarding table, router 1d must determine which gateway it should forward packets towards for dest *x*
 - ❖ this is also job of inter-AS routing protocol!
- *hot potato routing policy: send* packet towards closest of two routers.

