

Software Defined Networking

Richard T. B. Ma

School of Computing

National University of Singapore

Material from: Scott Shenker (UC Berkeley), Nick McKeown (Stanford), Jennifer Rexford (Princeton)

CS 3103: Compute Networks and Protocols

An Intern Position

- ❑ The IT department of A*STAR's Bioinformatics Institute (BII) is seeking an intern for training in network administration and programming. The focus of the internship is on software-defined networks and Open Network Install Environment-related work so it will be interesting for those students who are keen with new technologies. Interested students please submit the followings
 - - detailed CV
 - - copy of NRIC or passport (if foreigner)
 - - copy of educational certificates & transcripts
 - - period available for internship
- ❑ To Ms Betty Kee
 - Senior Admin Manager, Bioinformatics Institute (BII), 30 Biopolis Street, #07-01 Matrix, Singapore 138671
 - Main Tel: (65) 6478 8298
 - Email: bettyk@bii.a-star.edu.sg
 - <http://www.bii.a-star.edu.sg/>

Motivation

- ❑ What is Software Defined Networking?
 - Is it OpenFlow? or what is OpenFlow?
 - How about Network Virtualization?
 - How about separation of Control/Data planes?
 - What is the software part of SDN?

- ❑ Before knowing what it is, let us understand why we need it in the first place
 - What is wrong with the current Internet?


The Internet: A Remarkable Story

- ❑ Tremendous success
 - From research experiment to global infrastructure
- ❑ Brilliance of under-specifying
 - Network: best-effort packet delivery
 - Hosts: arbitrary applications
- ❑ Enables innovation in applications
 - Web, P2P, VoIP, social networks, virtual worlds
- ❑ But, change is easy only at the edge... ☹



Inside the Net: A Different Story


- ❑ Closed equipment
 - Software bundled with hardware
 - Vendor-specific interfaces
- ❑ Over specified
 - Slow protocol standardization
- ❑ Few people can innovate
 - Equipment vendors write the code
 - Long delays to introduce new features



Impacts performance, security, reliability, cost...

Networks are Hard to Manage

- ❑ Operating a network is expensive
 - More than half the cost of a network
 - Yet, operator error causes most outages
- ❑ Buggy software in the equipment
 - Routers with 20+ million lines of code
 - Cascading failures, vulnerabilities, etc.
- ❑ The network is “in the way”
 - Especially a problem in data centers
 - ... and home networks



Networks Vs. Other Systems

❑ Networks are hard to manage

- Computation and storage have been virtualized
 - Creating a more flexible and manageable infrastructure
- Networks are still notoriously hard to manage
 - Still heavily rely on network administrators

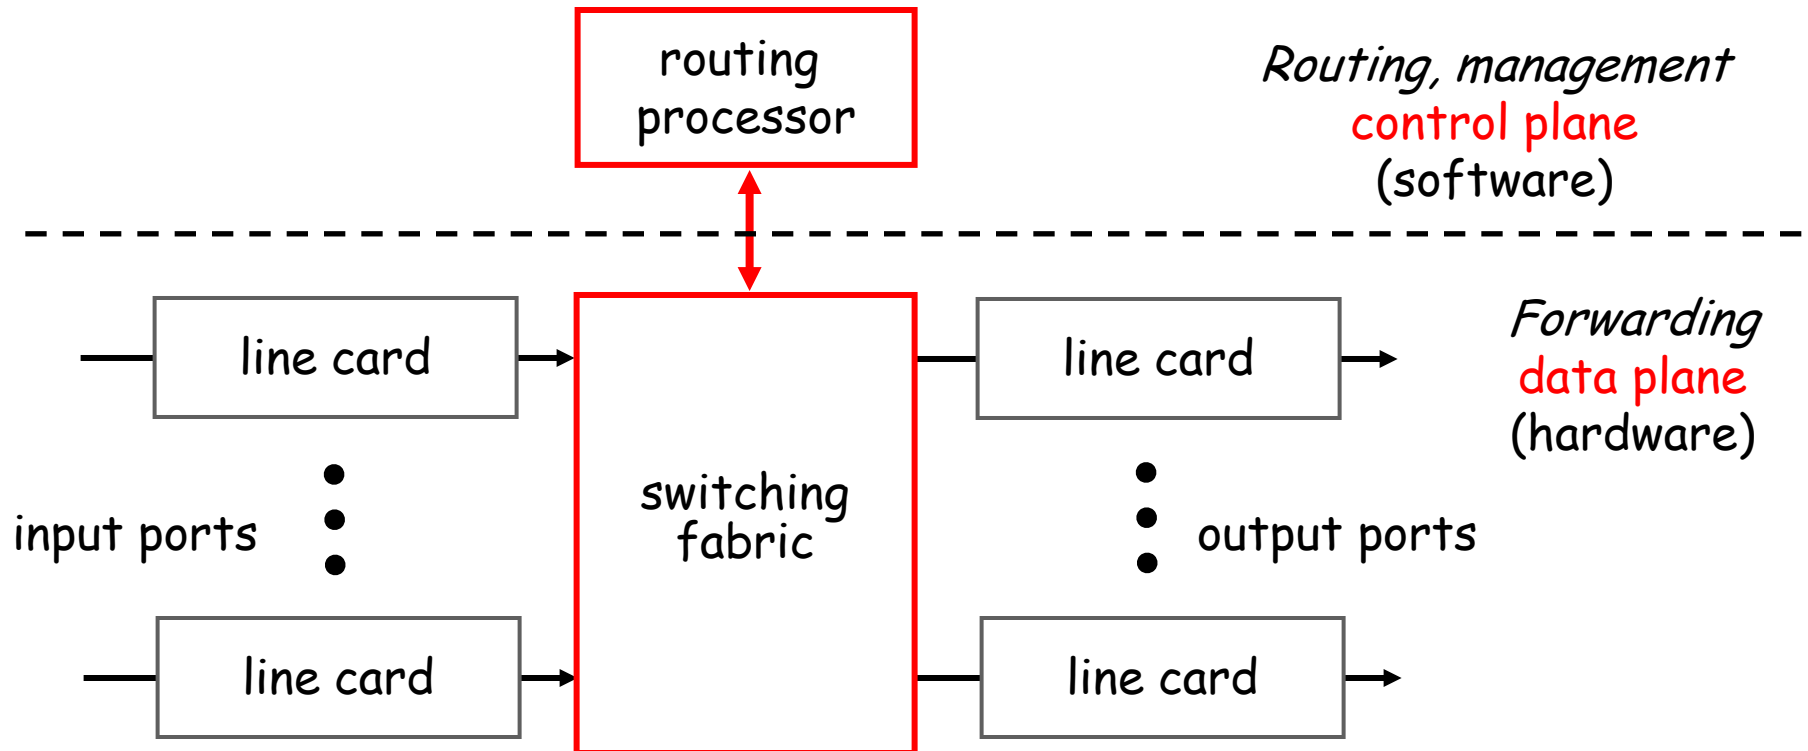
❑ Networks are hard to evolve

- Ongoing innovation in systems software
 - New programming languages, operating systems, etc.
- Networks are stuck in the past
 - Routing algorithms change very slowly
 - Network management extremely primitive

Complicated Router at the Core

two key router functions:

- ❑ run routing algorithms/protocol (RIP, OSPF, BGP)
- ❑ *forwarding* datagrams from incoming to outgoing link



Two Key Definitions

- **Data Plane:** processing and delivery of packets
 - Based on state in routers and endpoints
 - E.g., IP, TCP, Ethernet, etc.
 - Fast timescales (per-packet)

- **Control Plane:** establishing the state in routers
 - Determines how and where packets are forwarded
 - Routing, traffic engineering, firewall state, ...
 - Slow time-scales (per control event)

Networking as a Discipline

- ❑ Other fields in “systems”: OS, DB, DS, etc.
 - Teach basic principles
 - Are easily managed
 - Continue to evolve

- ❑ Networking:
 - Teach big bag of protocols
 - Notoriously difficult to manage
 - Evolves very slowly

A failure from an academic point of view

Why Does Networking Lag Behind?


- ❑ Networks used to be simple: Ethernet, IP, TCP....
- ❑ New **control** requirements led to great complexity
 - Isolation → VLANs, ACLs
 - Traffic engineering → MPLS
 - Packet processing → Firewalls, NATs
 - Payload analysis → Deep packet inspection
- ❑ Mechanisms designed and deployed independently
 - Complicated “control plane” design, primitive functionality
 - Stark contrast to the elegantly modular “data plane”
- ❑ Ability to master complexity → a curse
 - Extract simplicity is needed to build a discipline

A Good Example: Programming

- ❑ Machine languages: no abstractions
 - Mastering complexity was crucial
- ❑ Higher-level languages: OS and other abstractions
 - File system, virtual memory, abstract data types, ...
- ❑ Modern languages: even more abstractions
 - Object orientation, garbage collection,...

Abstractions key to extracting simplicity

Key to the Internet's Success



- ❑ Hourglass IP model
- ❑ Layered service abstractions (why is this important?)
 - Decompose delivery into fundamental components
 - Independent, compatible innovation at each layer
- ❑ Only for network edges

What have we learned so far?

- ❑ Layers are great abstractions
 - Layers only deal with the data plane
 - No powerful control plane abstractions!

“Modularity based on abstraction is the way things get done” - Barbara Liskov

Abstractions → Interfaces → Modularity

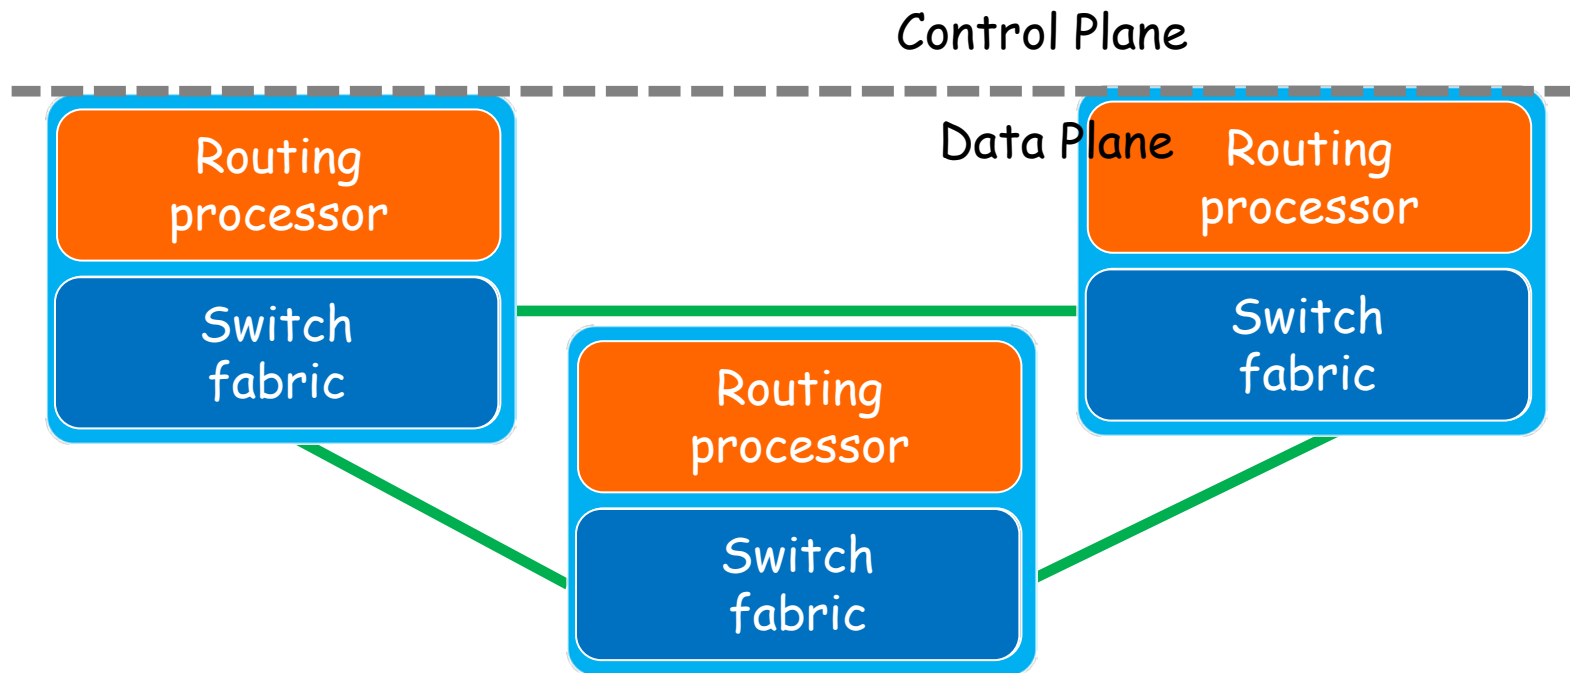
- ❑ How do we find control plane abstractions?
 - first define problem
 - and then decompose it

The network control problem

- ❑ Compute the configuration of each physical devices, e.g., forwarding tables
- ❑ Operate without communication guarantees
- ❑ Operate within given network-level protocol, e.g., RIP, OSPF.

Only people who love complexity would find this a reasonable request!

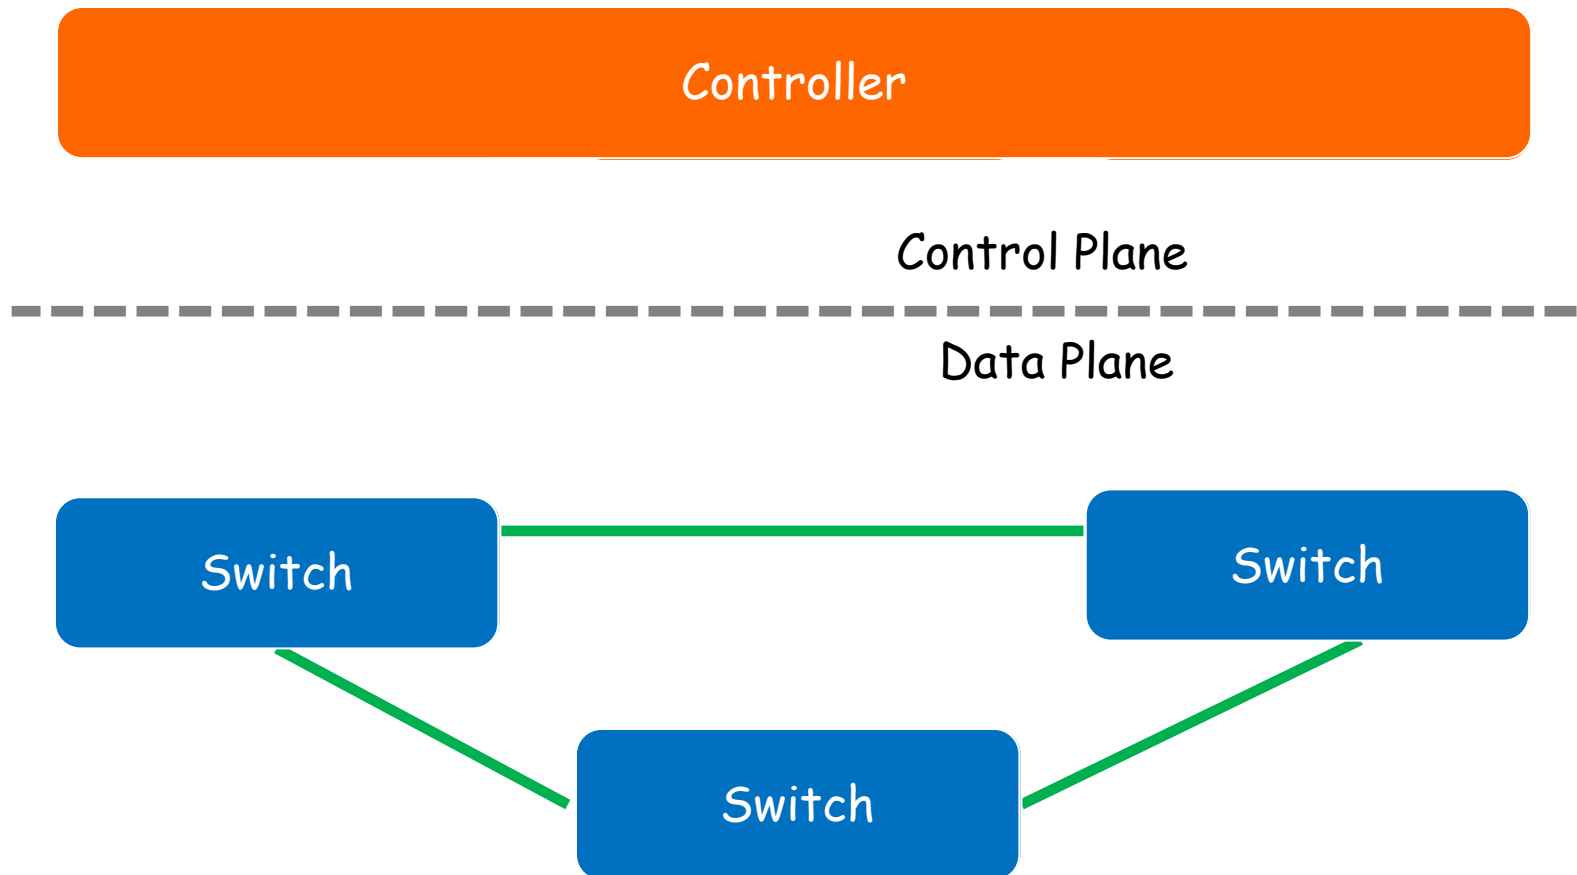
Separation of Control/Data Plane



Benefits of Separation

- ❑ Independent evolution and development
 - The software control of the network can evolve independently of the hardware.
- ❑ Control from high-level software program
 - Control behavior using higher-order programs
 - Debug/check behavior more easily

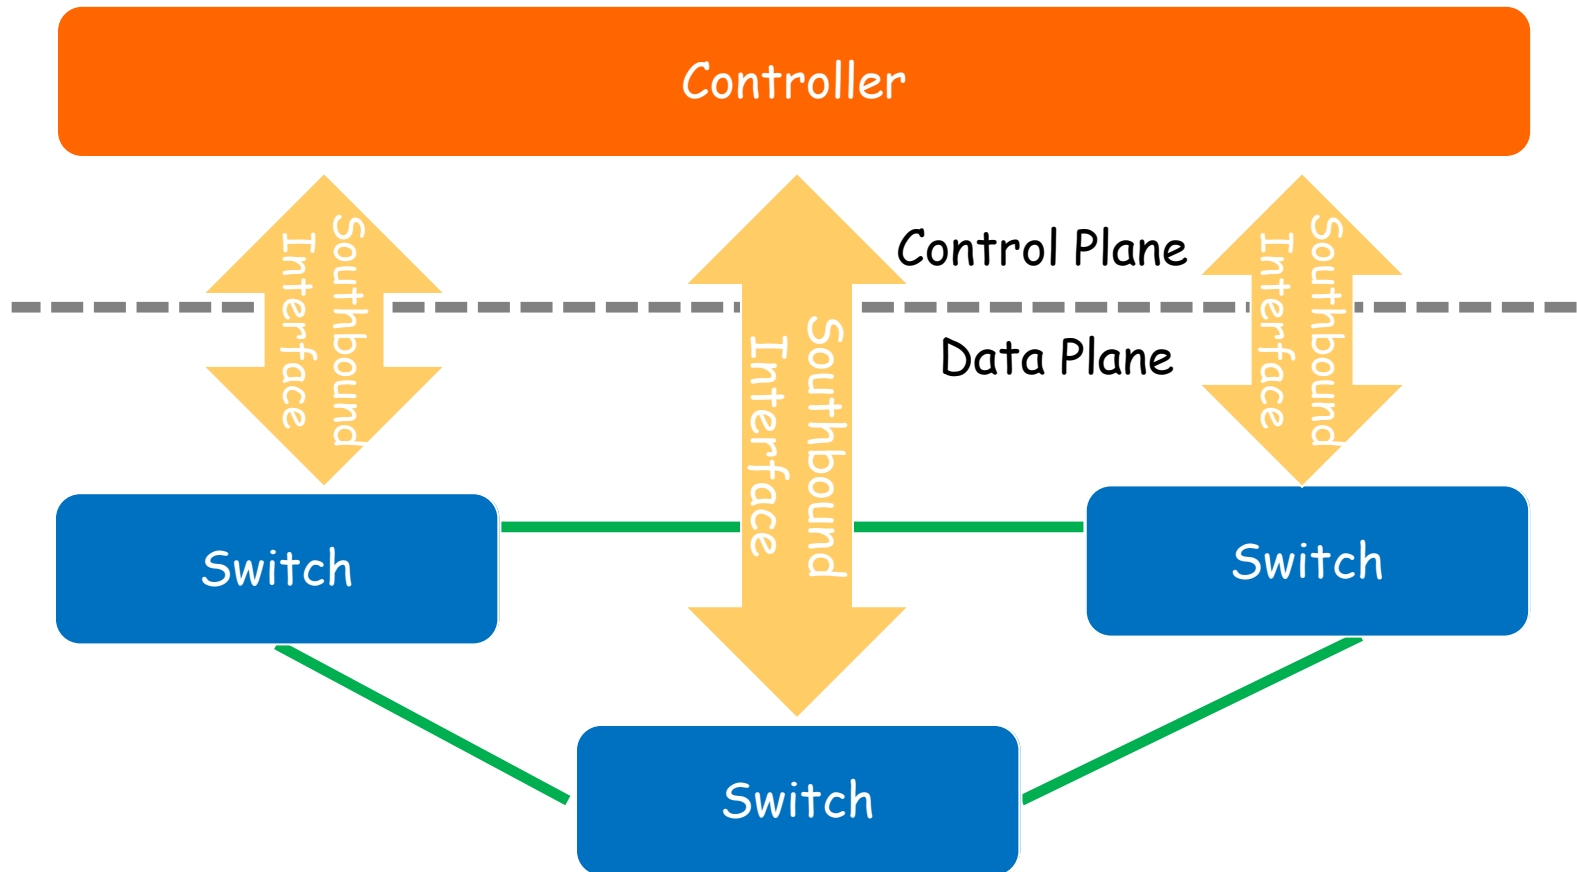
Logically Centralized Control



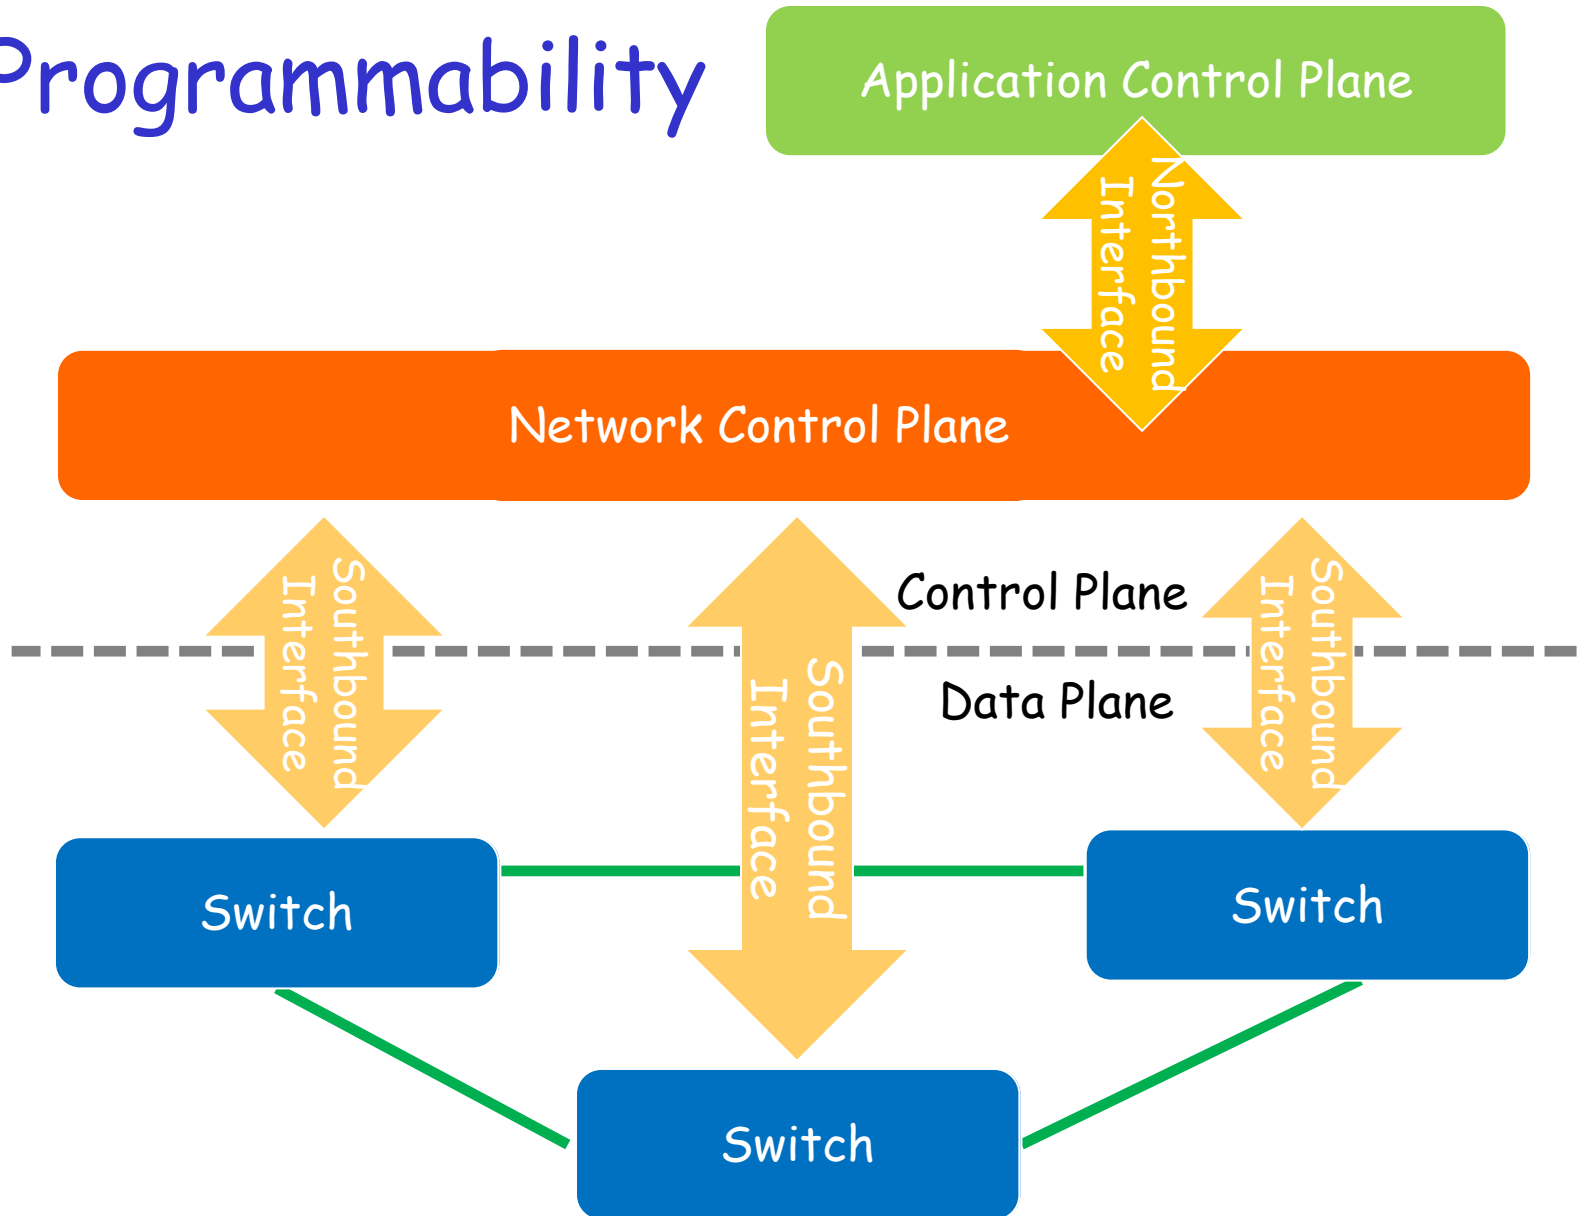
Benefits of Centralization

- ❑ Centralized decisions are easier to make
 - E.g., OSPF (RFC 2328) 244 pages
 - Distributed system part (builds consistent network 100 pages)
 - Routing algorithm (Dijkstra's algorithm 4 pages)
- ❑ Logically vs. physically centralized
 - Issues with of a physically centralized controller?
 - How to implement a logically centralized one?

Open Interfaces



Programmability




Benefits of Open Interfaces and Programmability

- ❑ Enable competitive technologies
 - Independent developments
 - Rapid innovation and fast evolution
 - Cheap and better networks
- ❑ Make network management much easier
 - Management goals are expressed as policies
 - New control/services for network providers
 - Detailed configuration are done by controller

A Quick Summary

- ❑ Principles of Software Defined Networking
 - Separation of Control Plane and Data Plane
 - Logically Centralized Control
 - Open Interfaces
 - Programmability
- ❑ All these principles use abstractions to modularize the network control problem.
- ❑ A nice analogy from Prof. Nick McKeown

Mainframes




Vertically integrated
Closed, proprietary
Slow innovation
Small industry



— Open Interface —



— Open Interface —




Horizontal
Open interfaces
Rapid innovation
Huge industry

Routers/Switches



Vertically integrated
Closed, proprietary
Slow innovation




— Open Interface —



— Open Interface —

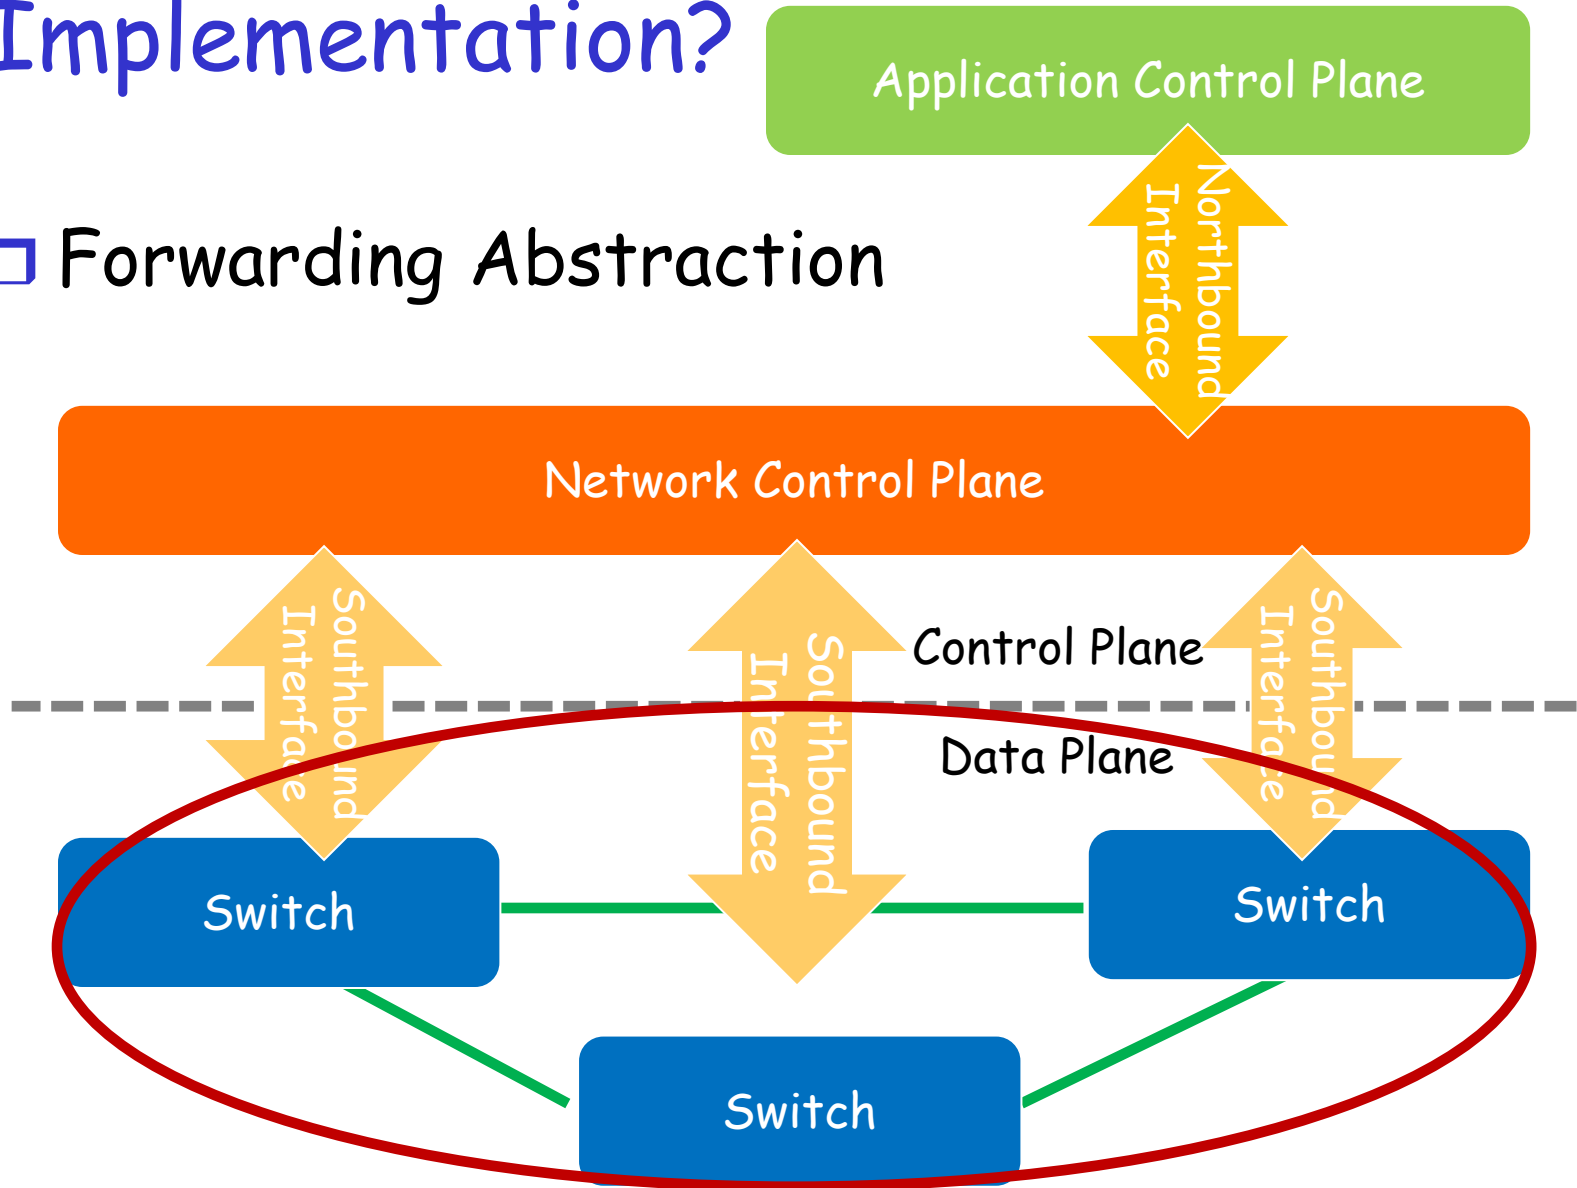


Horizontal
Open interfaces
Rapid innovation



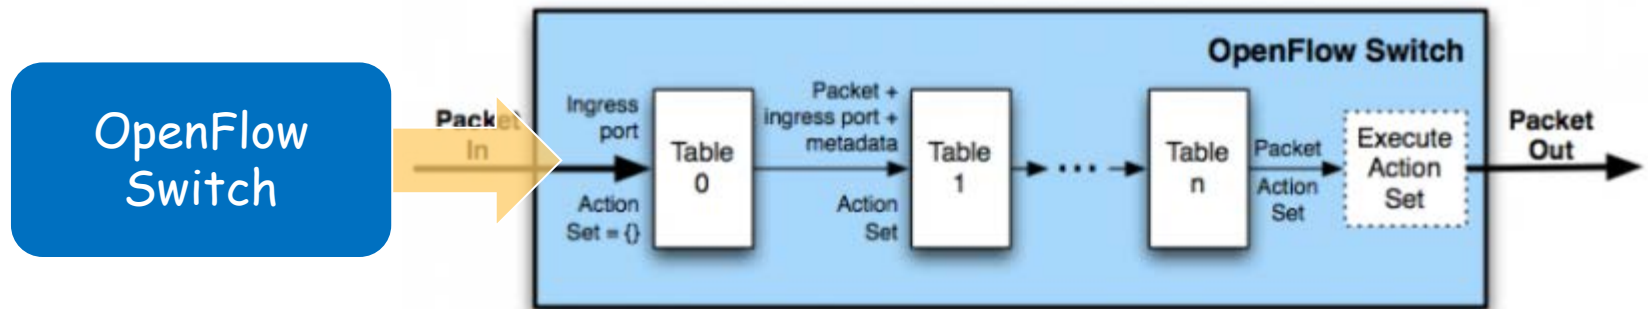
Implementation?

□ Forwarding Abstraction



OpenFlow Protocol and Switch

- ❑ OpenFlow switches have flow tables
 - Each entry is a (Match, Actions) primitive
 - Match: packets are compared with flow entries
 - Actions: depending on whether a match is found
 - If no match, traffic is sent to controller



- ❑ communicate with controller via secure channel

Match Fields

Ingress Port	Ethernet			VLAN		IP				TCP/UDP	
	SA	DA	Type	ID	Pri	SA	DA	Proto	TOS	Src	Dst

Flow Table

Classifier	Action	Statistics
Classifier	Action	Statistics
:	:	:
Classifier	Action	Statistics

of packets
of bytes

Actions

Forward	Physical Port	
	Virtual Port	ALL
		CONTROLLER
		LOCAL
		TABLE
		IN_PORT
DROP		
Forward	Virtual Port	NORMAL
		FLOOD
Enqueue		
Modify Field		

Mandatory Action

Optional Action

OpenFlow 1.0 primitives

Flow Actions

❑ Forward

- ALL: all but not including the incoming interface
- CONTROLLER: send the packet to the controller
- IN_PORT: send packet back to its incoming port

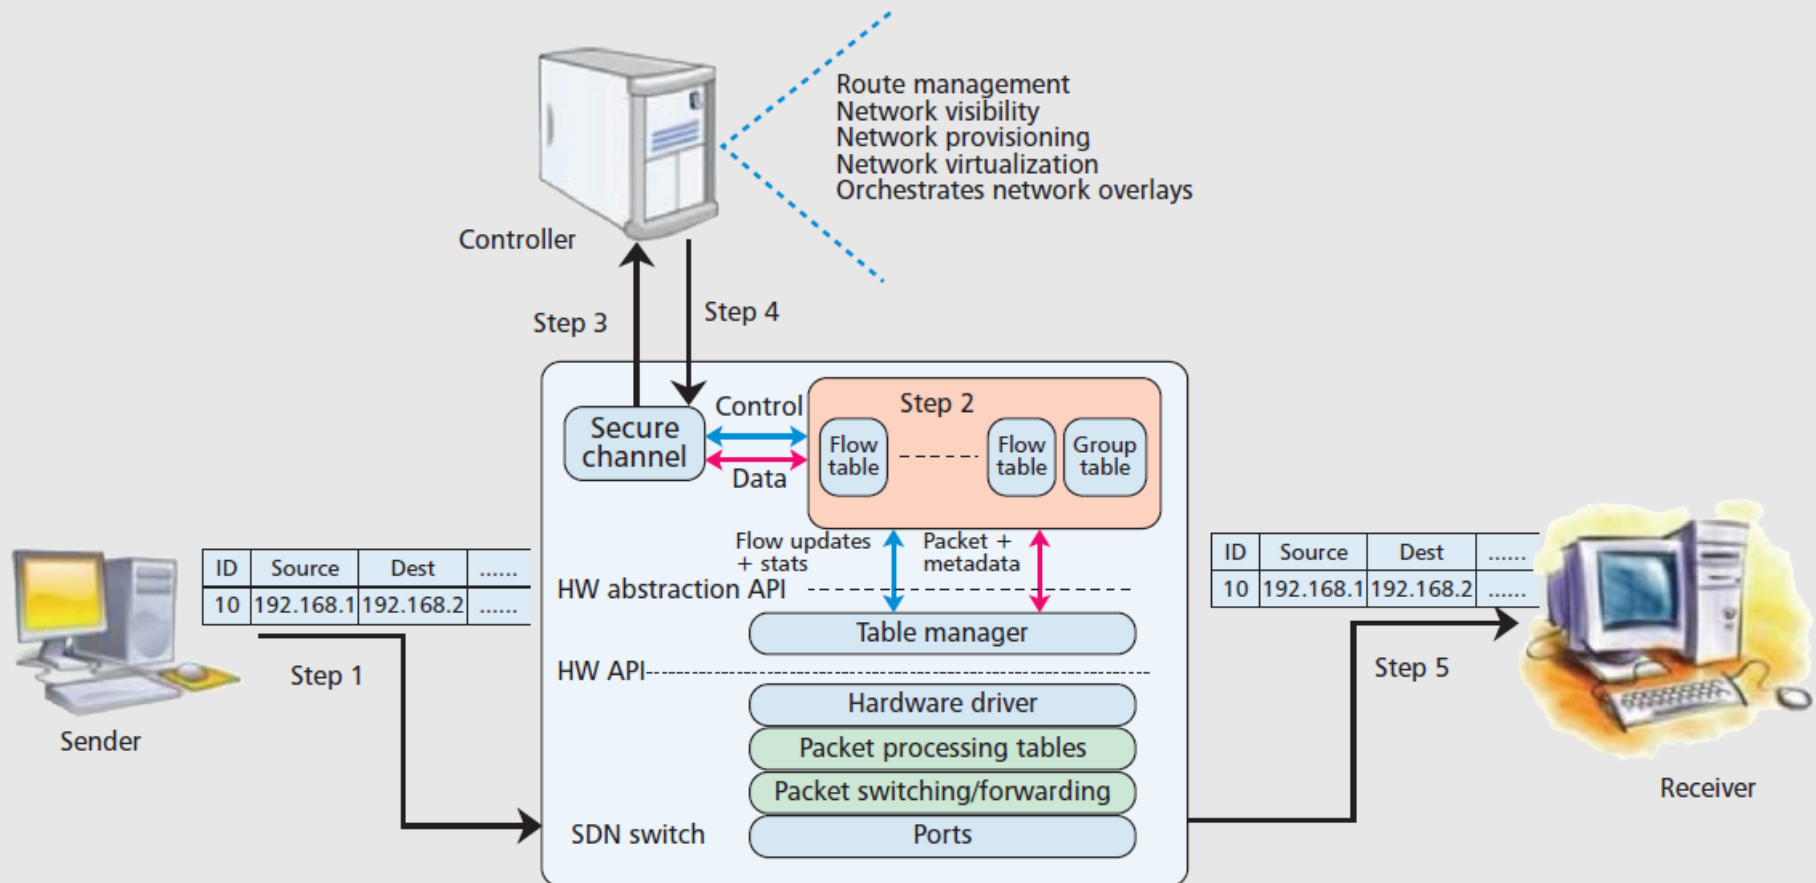
❑ Drop

- a flow-entry with no specified action indicates that all matching packets should be dropped

Flow Actions

- ❑ Modify: option to modify packet header values in the packet
 - set VLAN ID, priority, etc.
 - set destination IP address
- ❑ Enqueue: send the packet through a queue attached to a port
- ❑ More details from OpenFlow switch specification:
<https://www.opennetworking.org/images/stories/downloads/sdn-resources/onf-specifications/openflow/openflow-spec-v1.4.0.pdf>

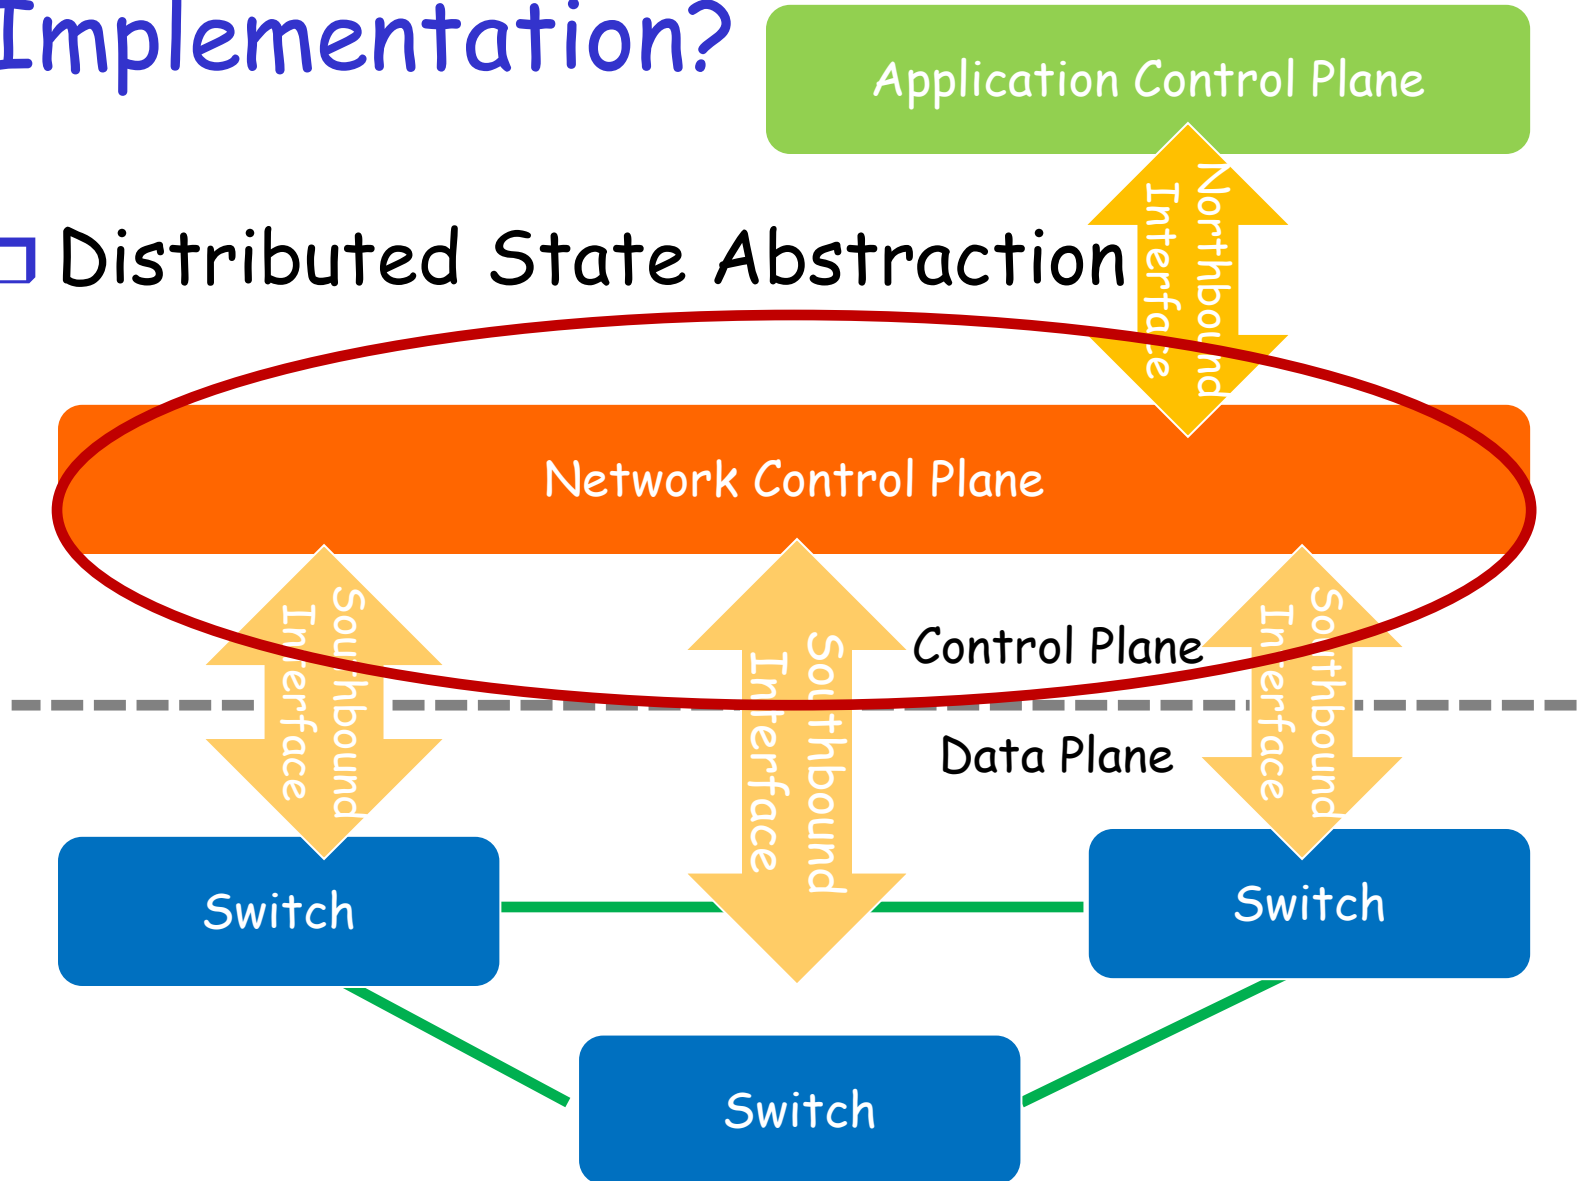
Operation of SDN (controller-switch)



S. Sezer et al. "Are we ready for SDN? Implementation Challenges for Software-Defined Networks", IEEE Communications Magazine, July 2013.

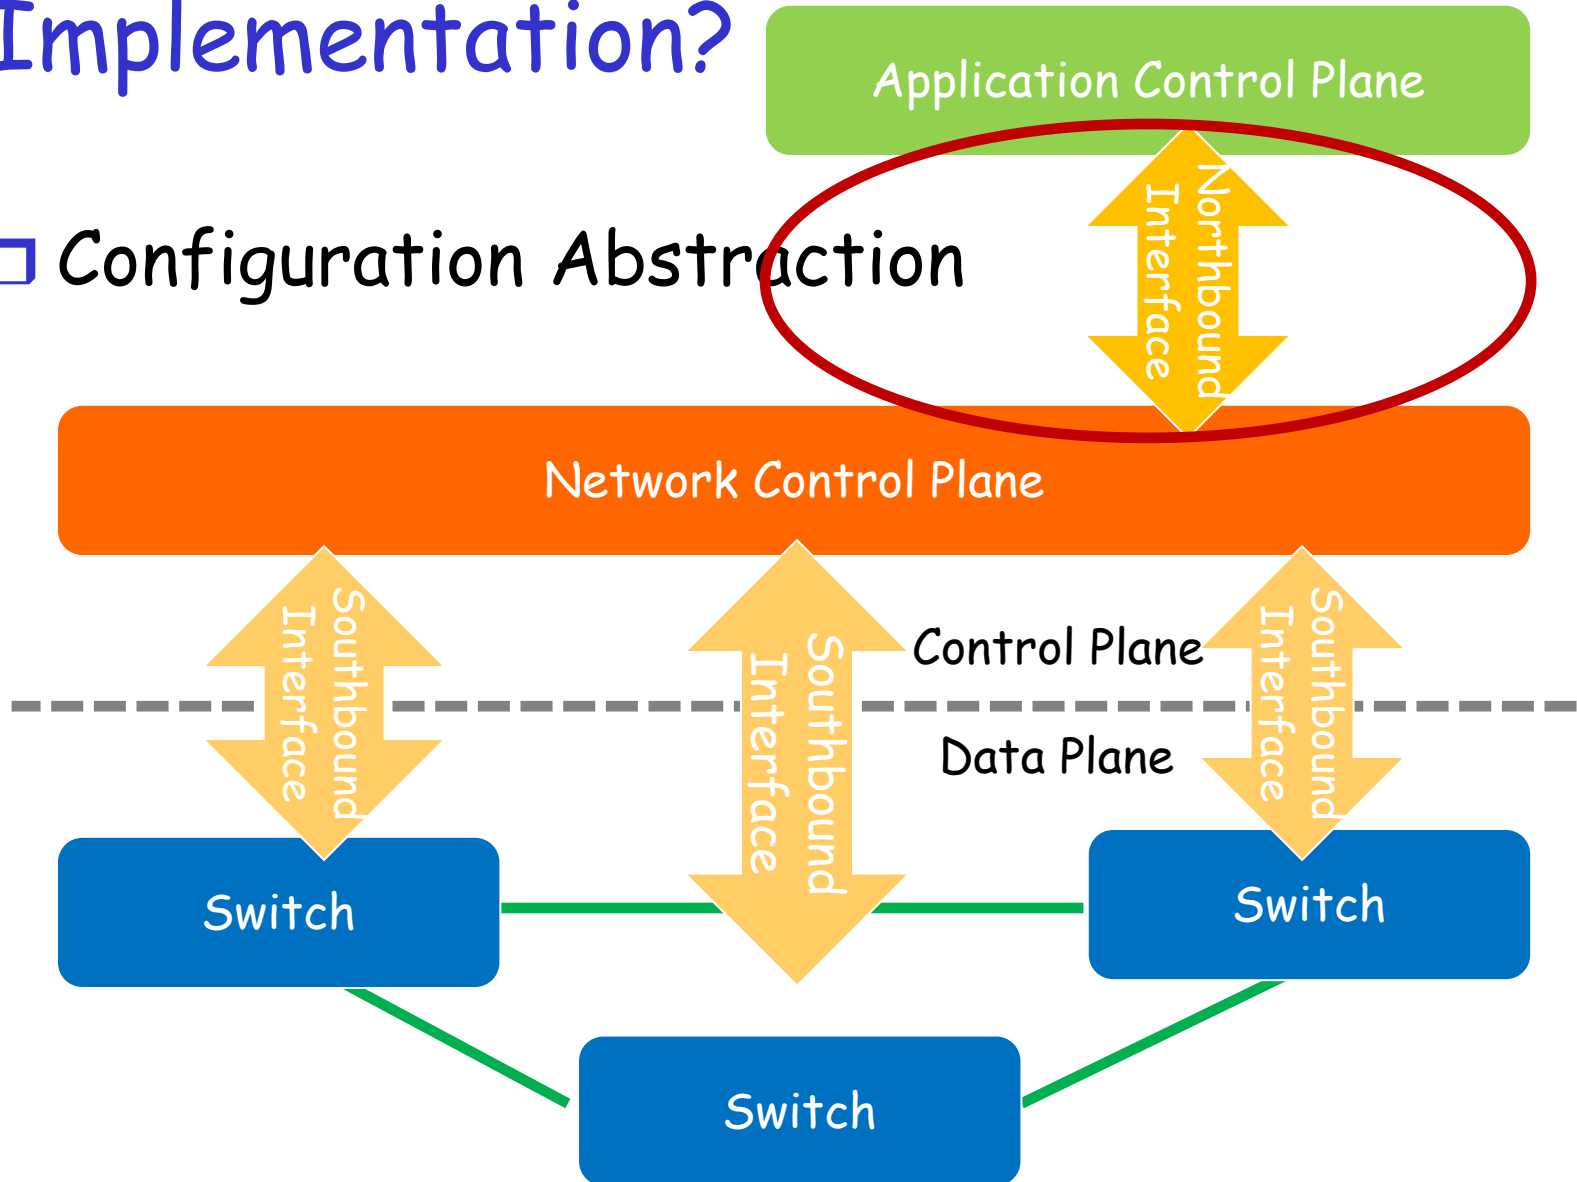
Implementation?

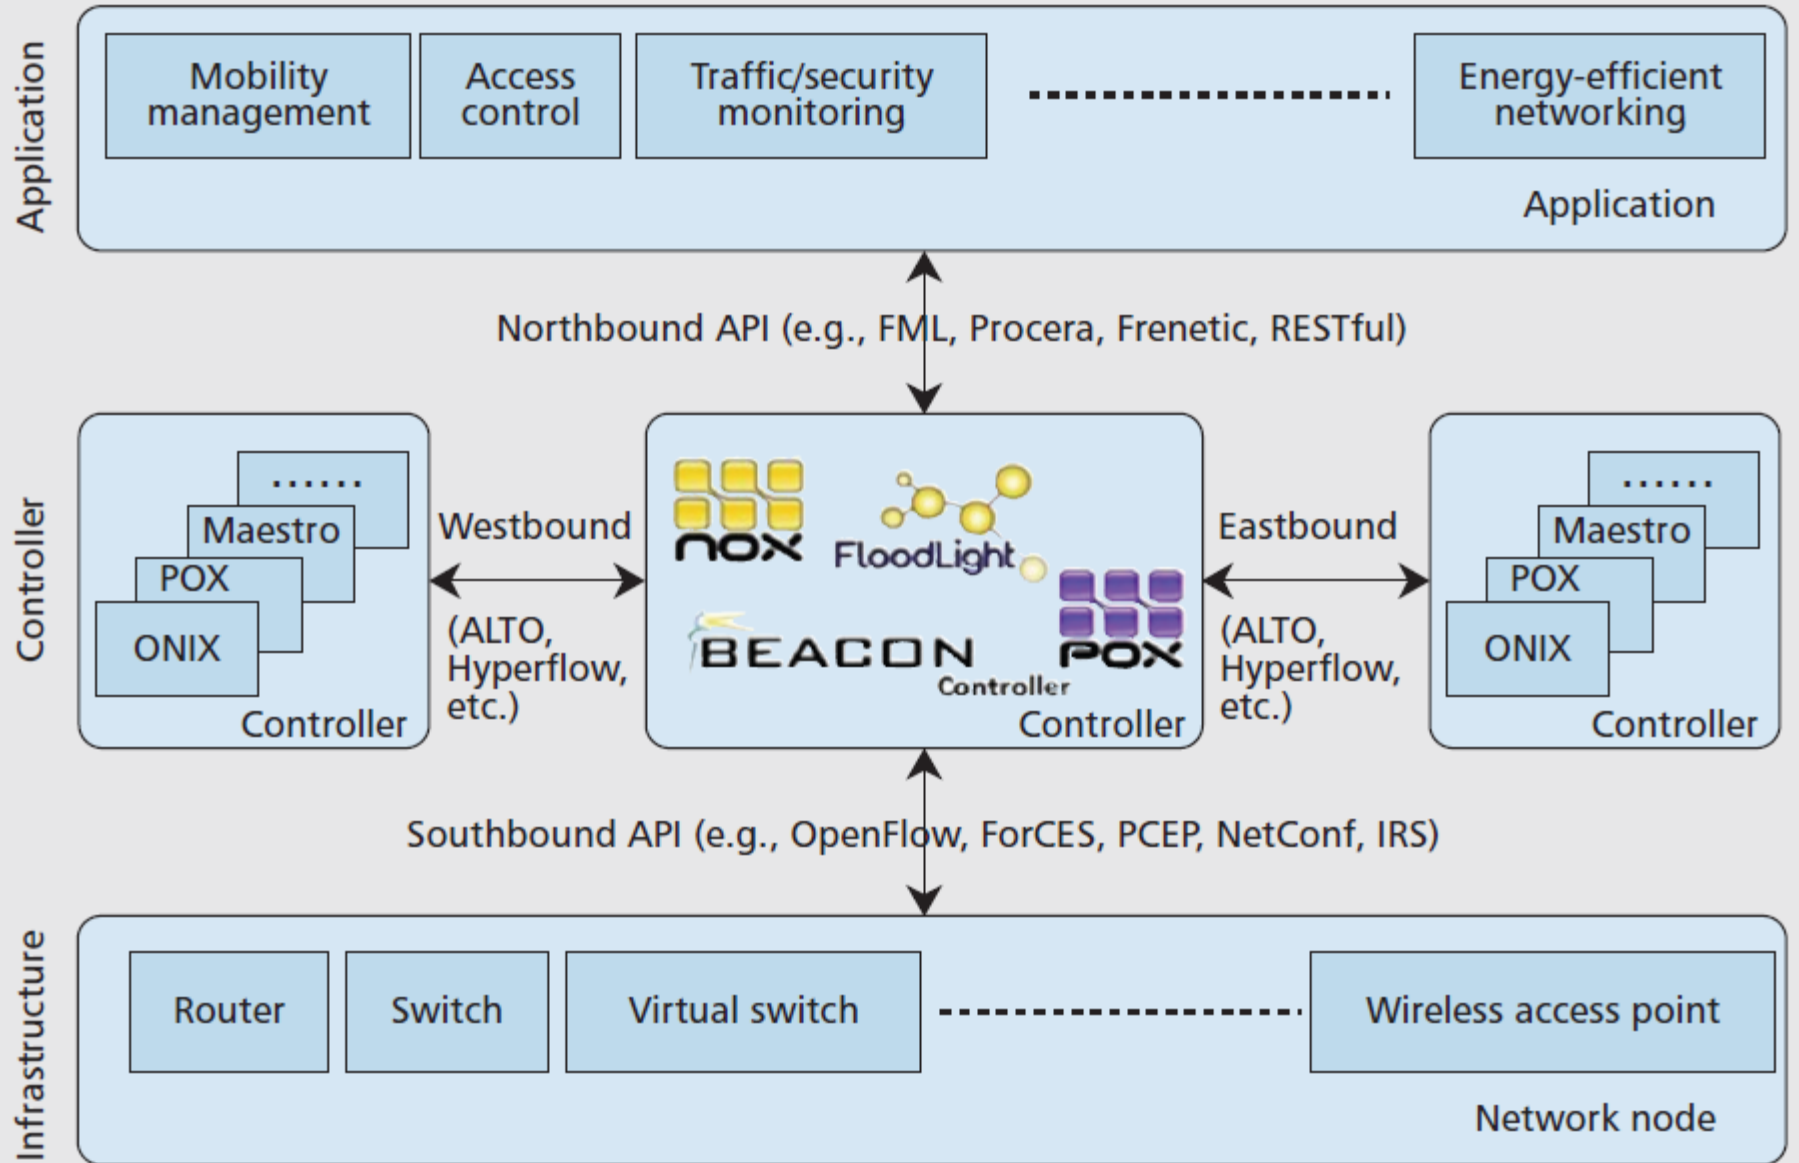
❑ Distributed State Abstraction



Implementation?

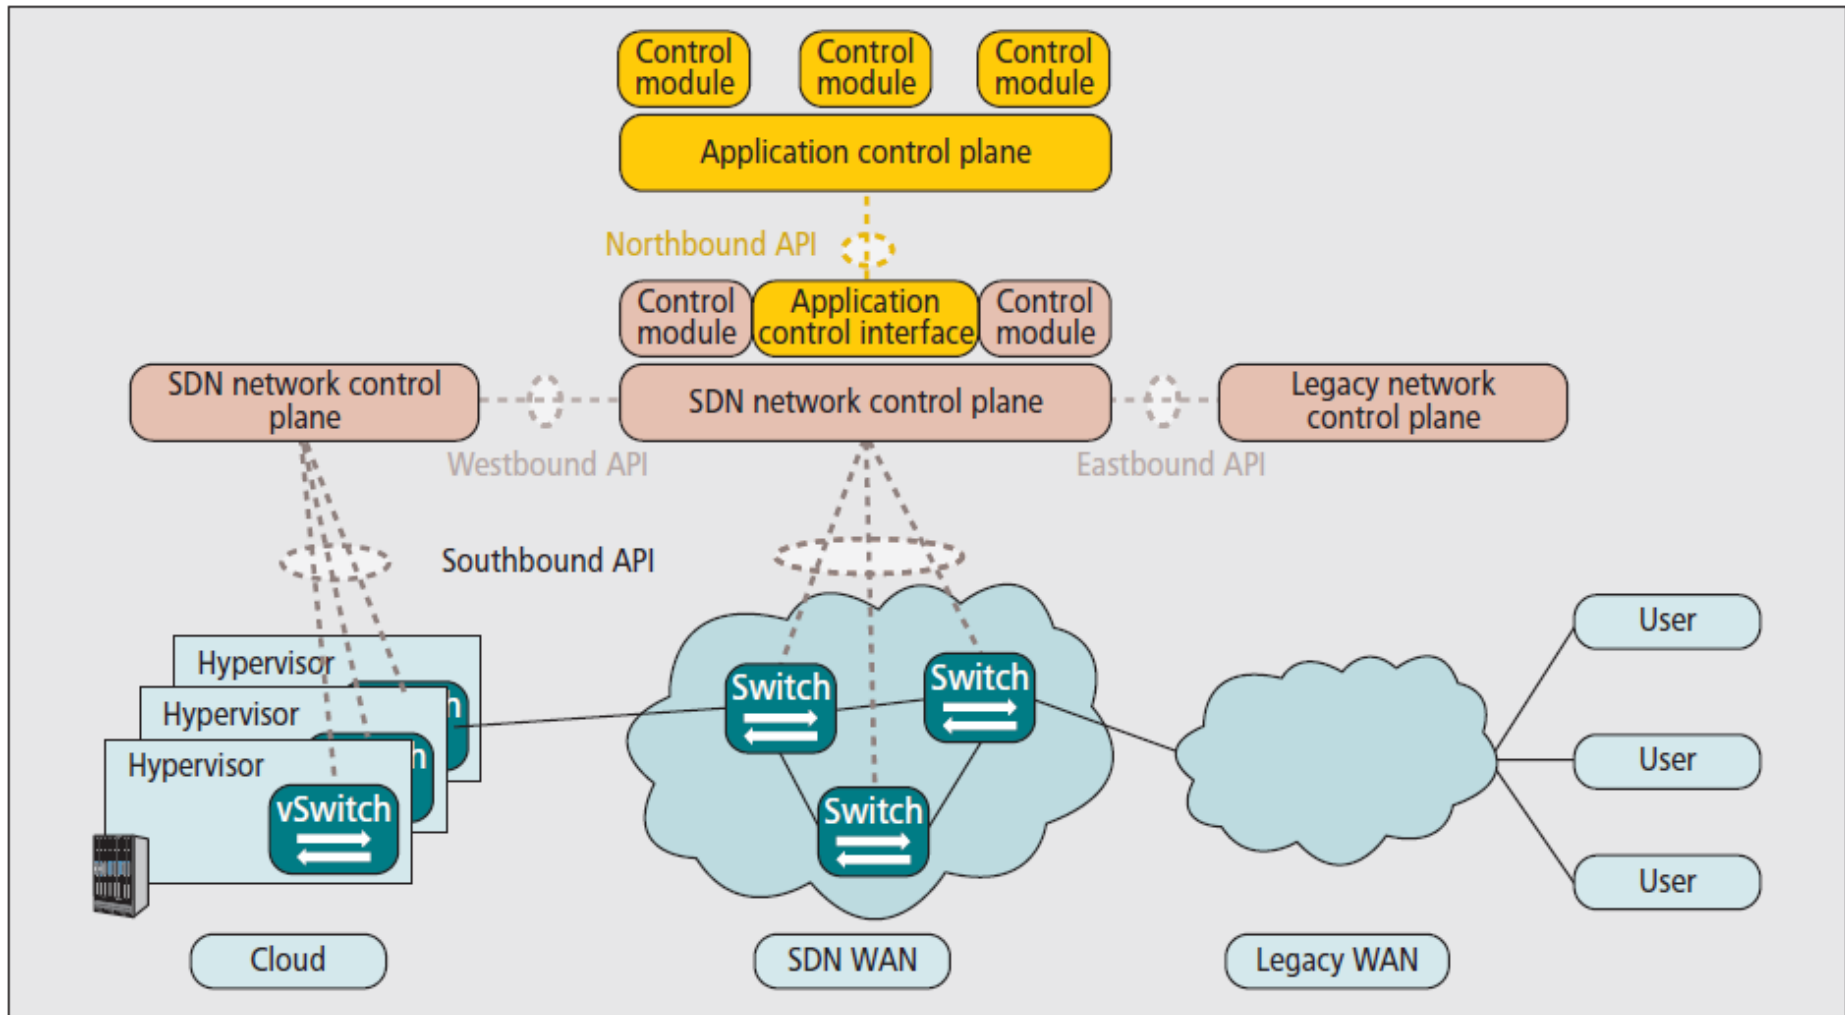
□ Configuration Abstraction





S. Sezer et al. "Are we ready for SDN? Implementation Challenges for Software-Defined Networks", IEEE Communications Magazine, July 2013.

Interfaces of a SDN



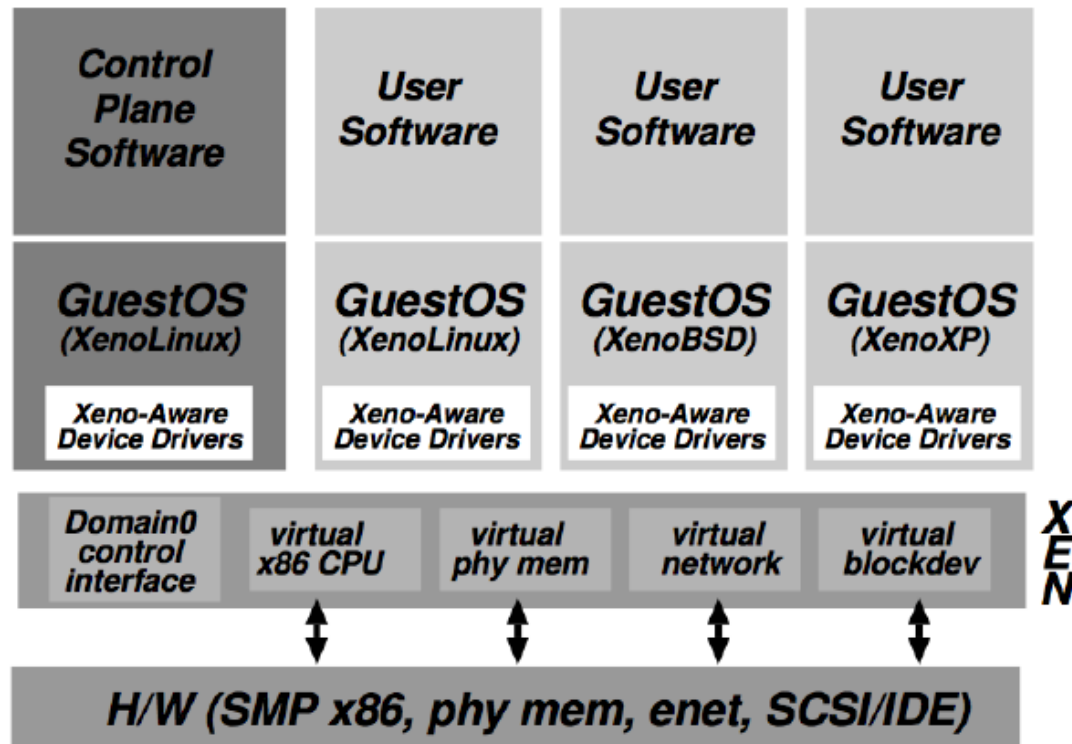
M. Jarschel et al. "Interfaces, attributes, and use cases: a compass for SDN", IEEE Communications Magazine, June 2014.

Network Virtualization (NV)

- ❑ What is virtualization?
 - Functionality of a single hardware is isolated for and used by multiple virtual instances
- ❑ Why do we want virtualization?
 - support multiple variations, easier for evolution
- ❑ What have been virtualized?
 - computing platform (virtual machines)
 - storage (virtual memory)
- ❑ How about networks?
 - not new, but very limited, e.g., VLAN and VPN


Example VM Environment: Xen

- ❑ Xen hosts multiple guest Oses
- ❑ Domain0 runs control software



Example Virtual Links: EGRE

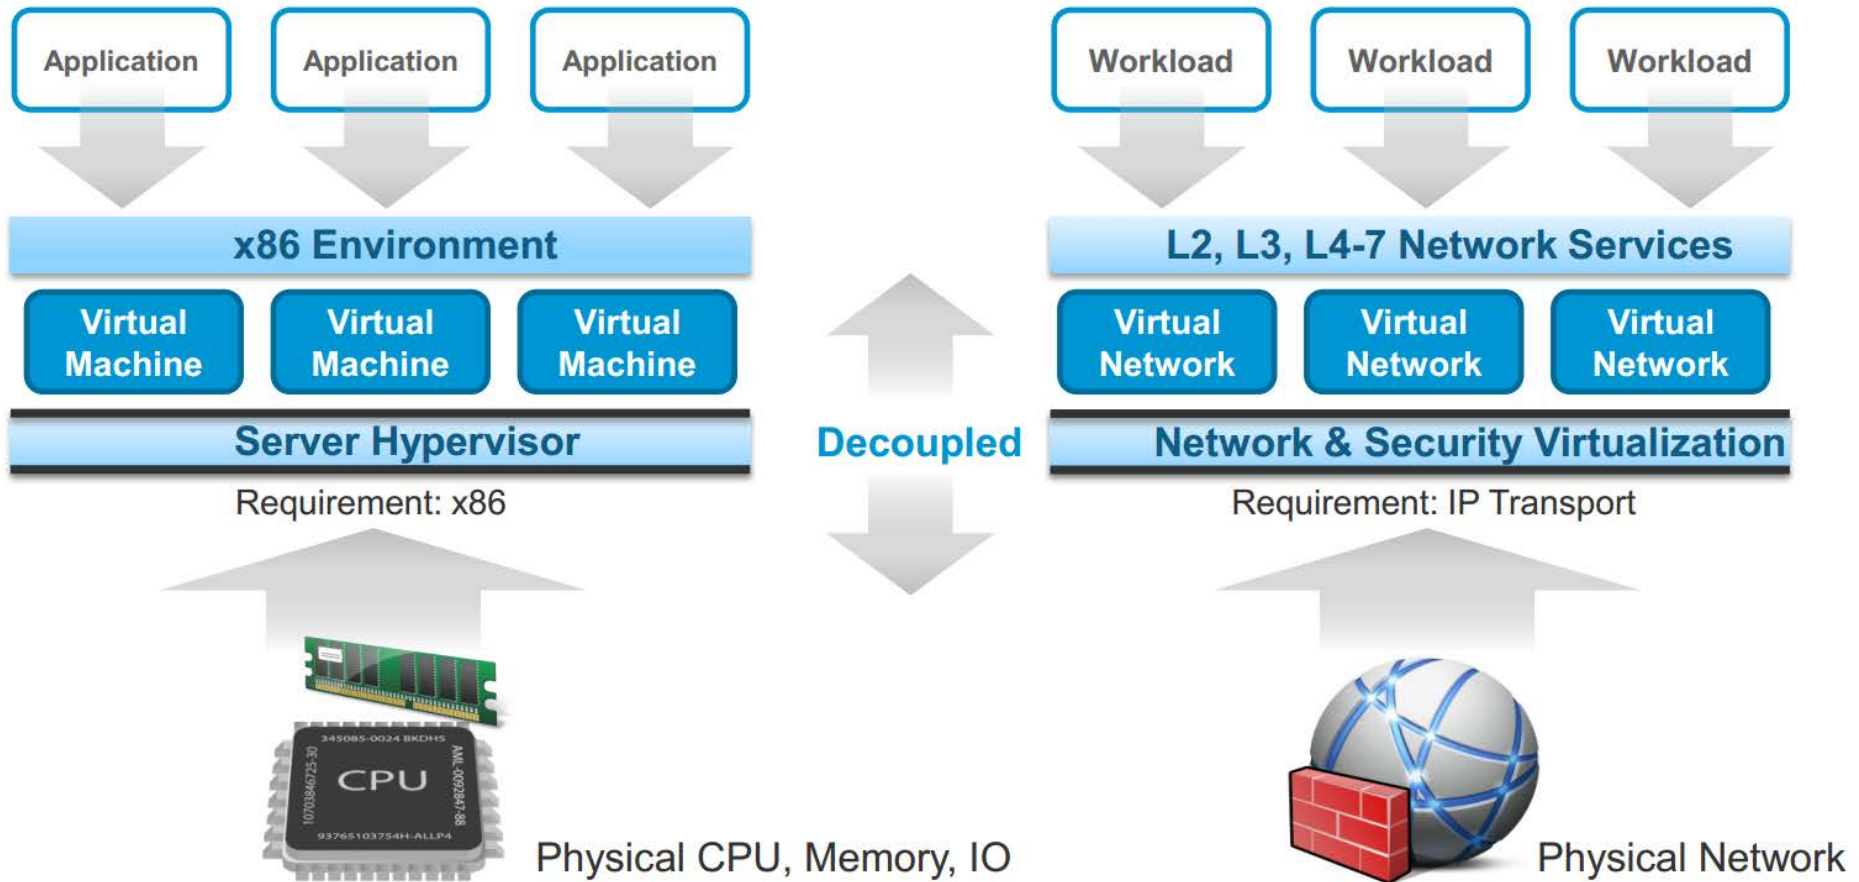
- ❑ Ethernet Generic Routing Encapsulation (EGRE): "Tunneling", Ethernet frames from virtual hosts are encapsulated in IP



Virtualizing Networks

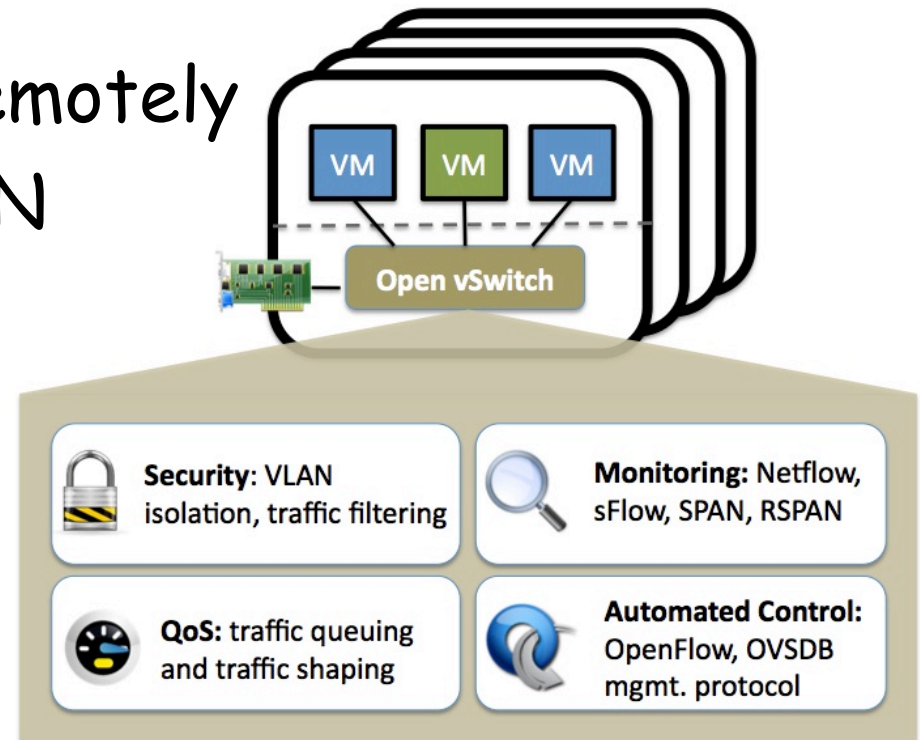
- ❑ network elements that can be virtualized
 - nodes: virtual machines
 - links: tunnels
 - storage
 - what else?
- ❑ abstraction of the whole physical network
 - support for multiple logical networks running on a common shared physical substrate
 - a container of network services
 - virtualization at different layers

Virtualizing Networks




Example Virtual Switch: OVS

- ❑ Open vSwitch is a virtual switch designed to enable massive network automation
- ❑ Can be configured remotely with OpenFlow, JSON



Example Flow Space Slicing

- ❑ divide the network into logical slices based on flow characteristics
- ❑ enforce traffic isolation among slices
- ❑ application: divide the production network into logical slices for experiments
 - each slice controls its own packet forwarding
- ❑ e.g., OpenFlow-based controller: FlowVisor



SDN != NV

- ❑ SDN alone does not abstract away details of physical network
- ❑ SDN is not required for NV, but SDN makes NV much easier
- ❑ NV is also a way to enable abstractions
- ❑ NV and SDN are orthogonal ...
 - SDN abstracts layers vertically
 - NV abstracts the same layer horizontally

Applications of SDN

❑ Killer app

- Network virtualization

❑ Other apps

- Datacenter and cloud computing
- WAN (Google B4)
- Software IXP
- Home API
- Wireless

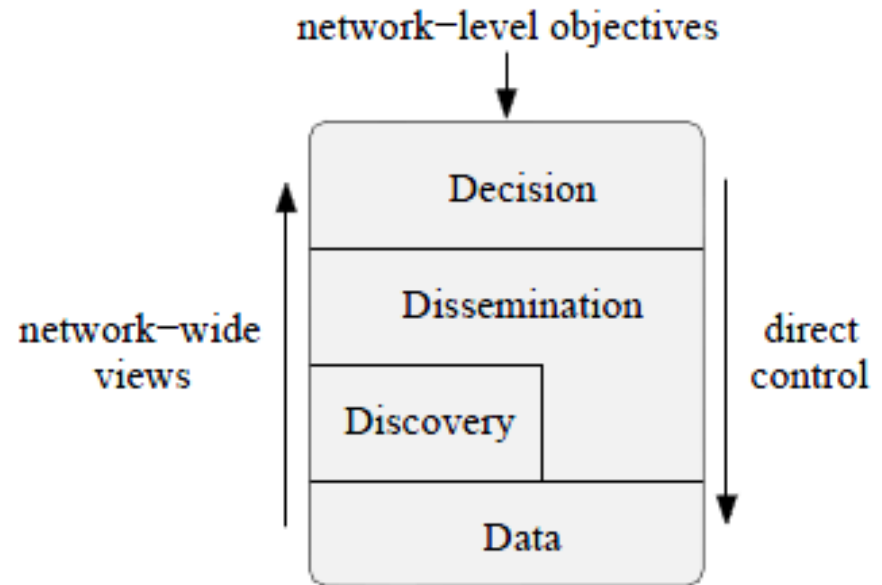
Related Work

❑ Principles of the 4D network architecture

- Network-level objectives
- Network-wide views
- Direct control

❑ Architecture


- Decision plane
- Dissemination plane
- Discovery plane
- Data plane



A. Greenberg et al. "A Clean Slate 4D Approach to Network Control and Management", *ACM SIGCOMM Computer Communication Review*, 35(5), 2005.

Related Work

- Routing Control Platform (RCP)
 - Logically centralized platform
 - Separate from the IP forwarding plane
 - Perform route selection using iBGP



M. Caesar et al. "Design and Implementation of a Routing Control Platform",
Proceedings of USENIX NSDI, 2005.



Related Work

❑ Ethane architecture

- Flow-based Ethernet switches with centralized controller

❑ Principles

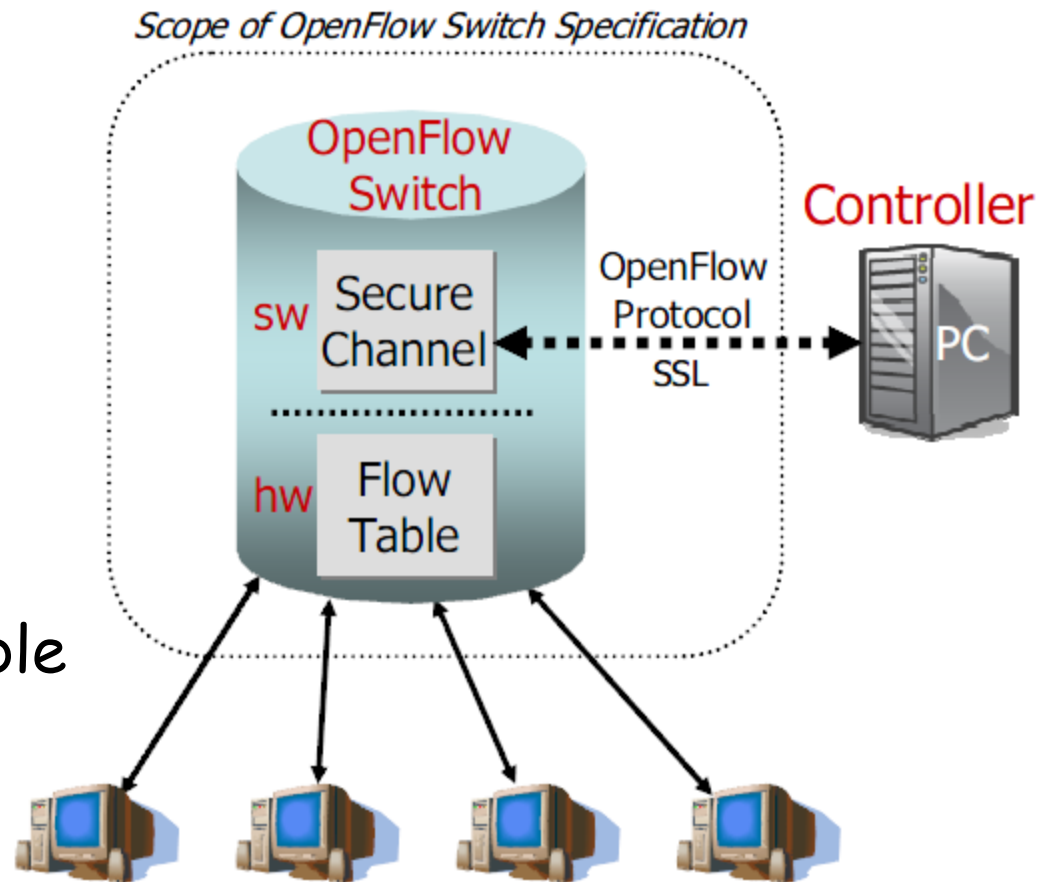
- Network governed by high-level policies
- Policy determine routing
- Binding between packets and its origin



M. Casado et al. "Ethane: Taking Control of the Enterprise", Proceedings of ACM SIGCOMM, 2007.

Related Work

- ❑ OpenFlow Switch
 - Ethernet switches with flow tables
 - $\langle \text{header}, \text{action} \rangle$
- ❑ OpenFlow Protocol
 - Enable programmable networks
 - Enable testing for new designs in production networks




N. McKeown et al. "OpenFlow: Enabling Innovation in Campus Networks", ACM SIGCOMM Computer Communication Review, 38(2), April, 2008.

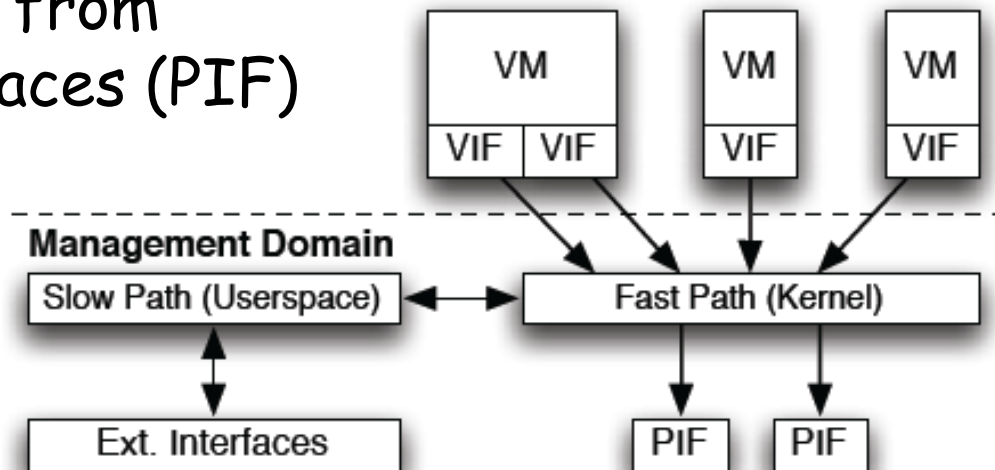
Related Work

❑ Open vSwitch

- Virtualize switch functionality in software
- Virtual interfaces (VIFs) are virtualized from physical interfaces (PIF)




Architecture of the Open vSwitch →



B. Pfaff et al. "Extending Networking into the Virtualization Layer", ACM Workshop on Hot Topics in Networks (HotNets), 2009.

Related Work



❑ Network OS

- Distributed systems
- NOX (~32K lines)
 - Developed side-by-side with OpenFlow
- Onix (~150K lines)
 - production-quality control platform

❑ Program Interface

- Event driven
- Network view
- High-level API

N. Gude et al. "NOX: Towards an Operating System for Networks", ACM SIGCOMM Computer Communication Review, 38(3), July, 2008.

T. Koponen et al. "Onix: A Distributed Control Platform for Large-scale Production Networks", USENIX OSDI, 2010.


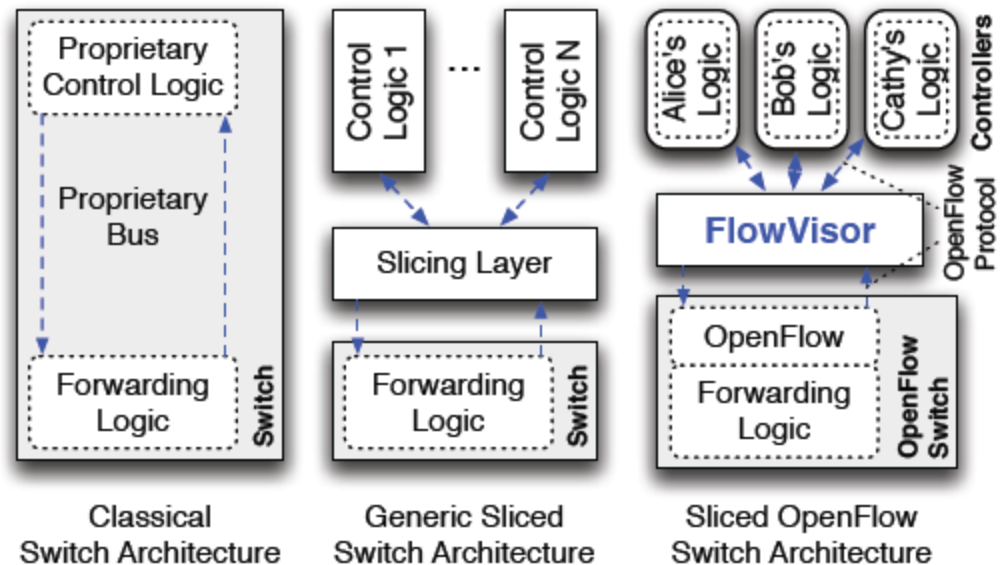
Related Work

❑ Virtualize network

- Flow space slicing
- A new slicing layer between forwarding and control planes

❑ FlowVisor

- Proxy for multiple OpenFlow controllers
- Used to experiment on production nets



R. Sherwood et al. "Can the Production Network Be the Testbed?", USENIX OSDI, 2010.

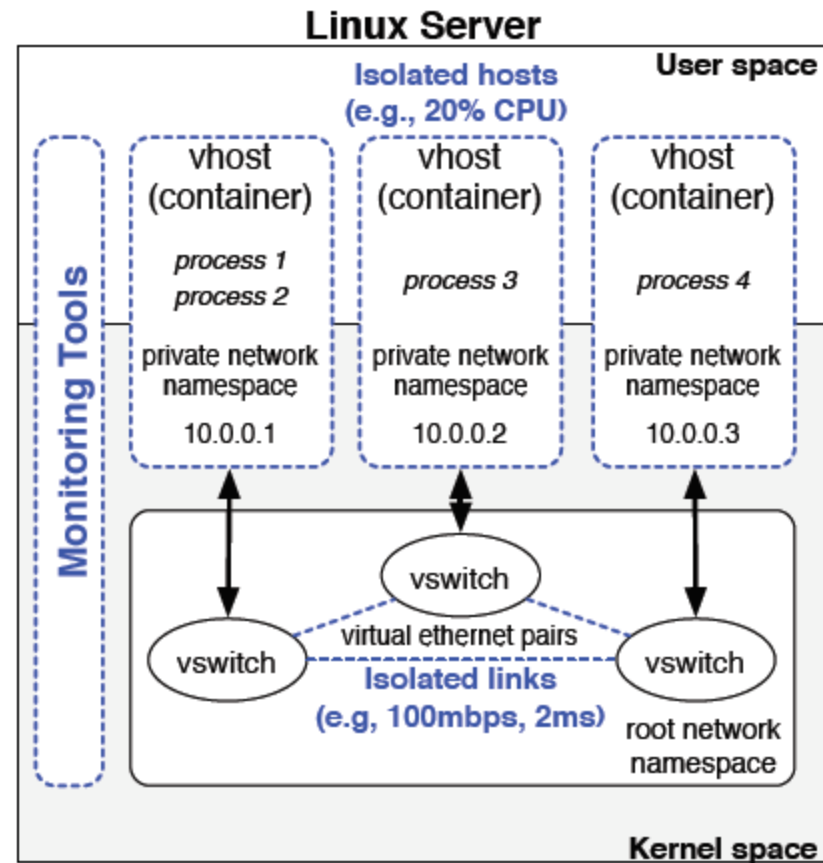

Related Work

❑ Mininet

- Emulation platform of network systems

❑ Design Features

- Using Linux Container mechanism
- Different namespaces for virtual hosts
- Virtual software switch implemented in Kernel

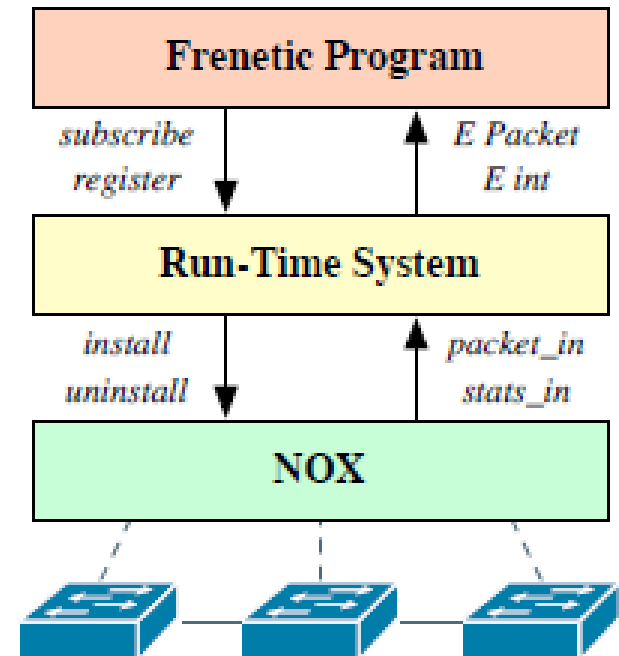


N. Handigol et al. "Reproducible Network Experiments Using Container-Based Emulation", ACM CoNEXT, 2012.

Related Work

□ Frenetic and Pyretic

- Control switches via policies
- North-bound APIs
- Frenetic: runtime system, functional reactive language
- Pyretic: Python & POX based, support multiple concurrent tasks via operation composition



N. Foster et al. "Frenetic: A Network Programming Language", ACM SIGPLAN International Conference on Functional Programming (ICFP), 2011.

C. Monsanto et al. "Composing Software-Defined Networks", Proceedings of USENIX NSDI, 2013.